

Deep Learning With TensorFlow

An Introduction To Artificial Neural Networks

By Brian Pugh

CMU Crash Course 1/28 2017

Goals

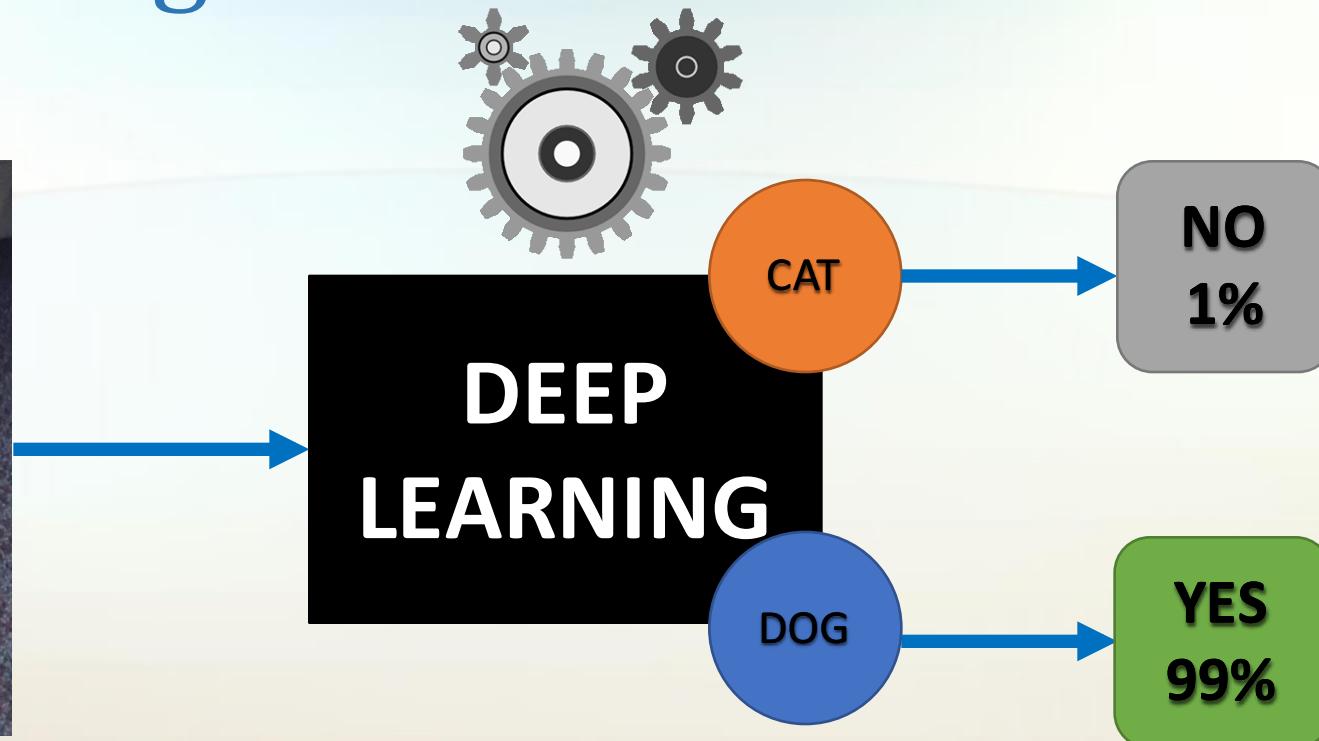
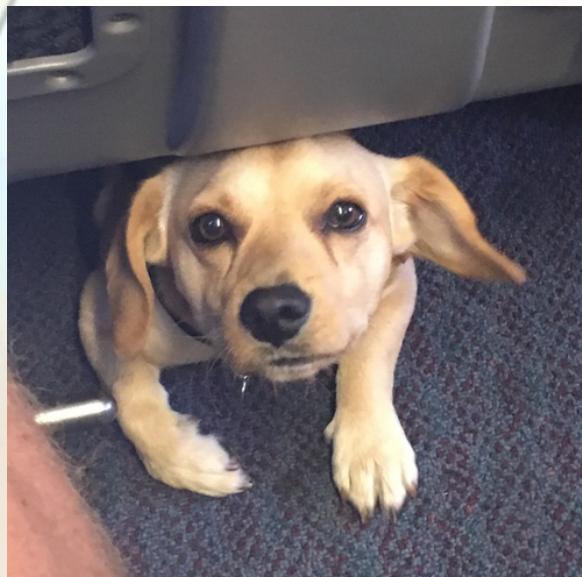
- What is Deep Learning?
- What is an Artificial Neural Network?
- A Basic Artificial Neural Network (Feed Forward Fully Connected)
- Implement a basic network for classifying handwritten digits

Deep Learning

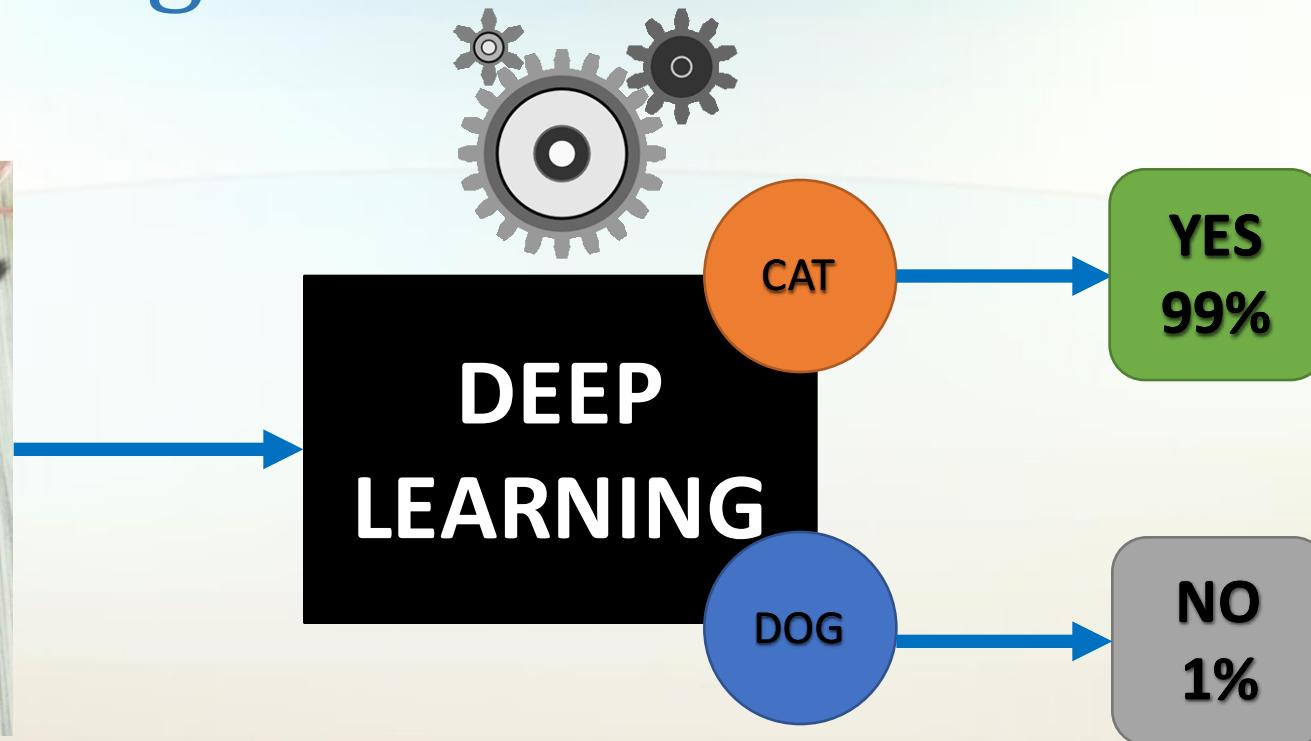
- A branch of Machine Learning
- Multiple levels of representation and abstraction
- One step closer to true “Artificial Intelligence”
- Typically refers to Artificial Neural Networks
- Externally can be thought of as a black box
- Maps inputs to outputs from rules it learns from training
- Training comes from known labeled input/output datasets.



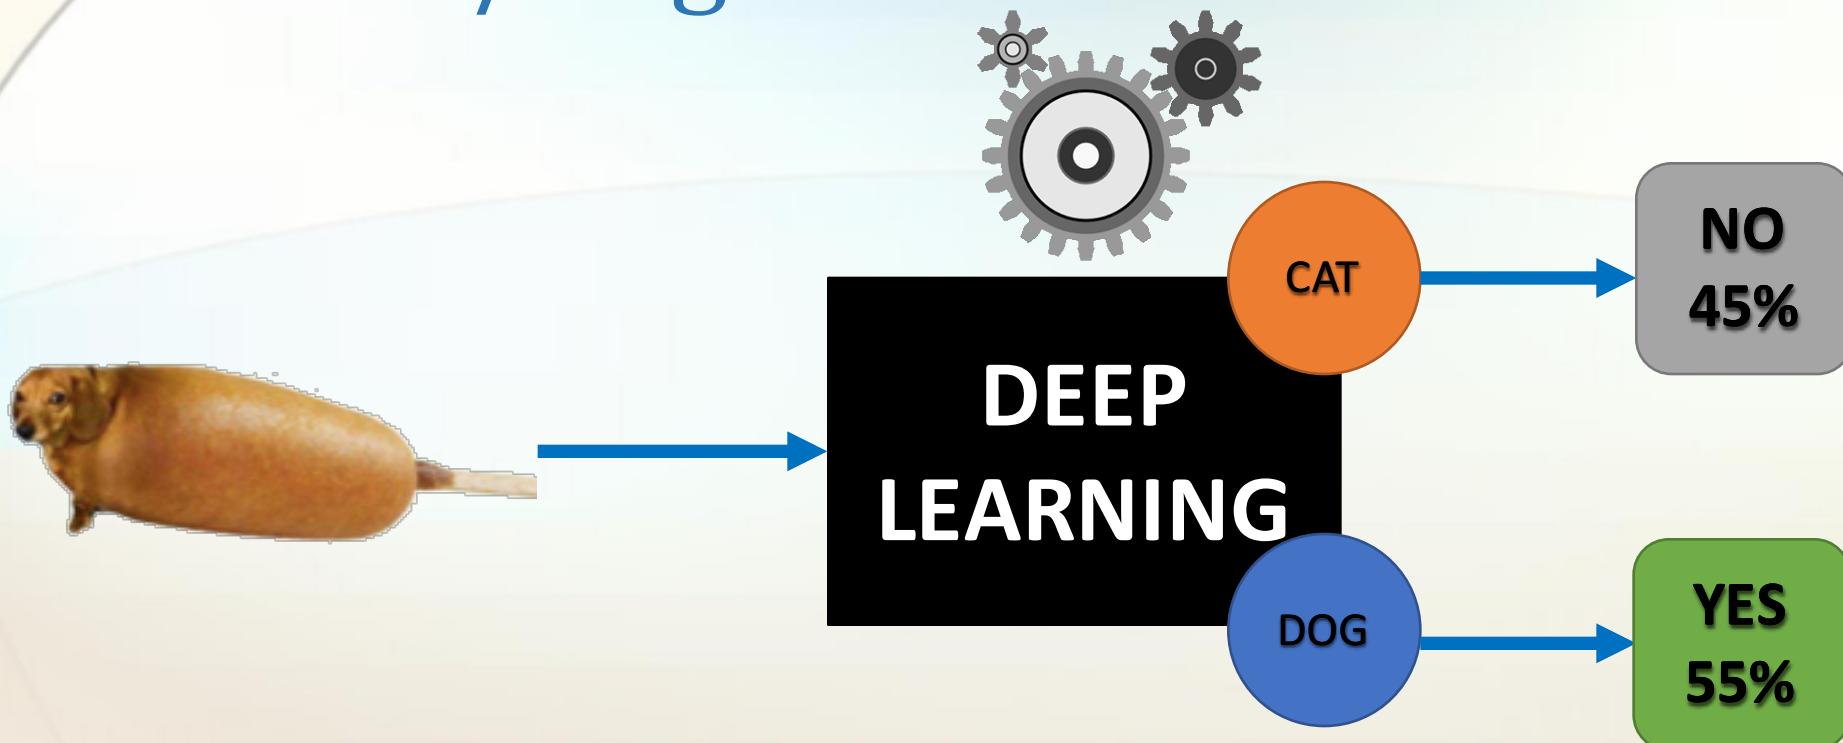
Cat/Dog Classifier



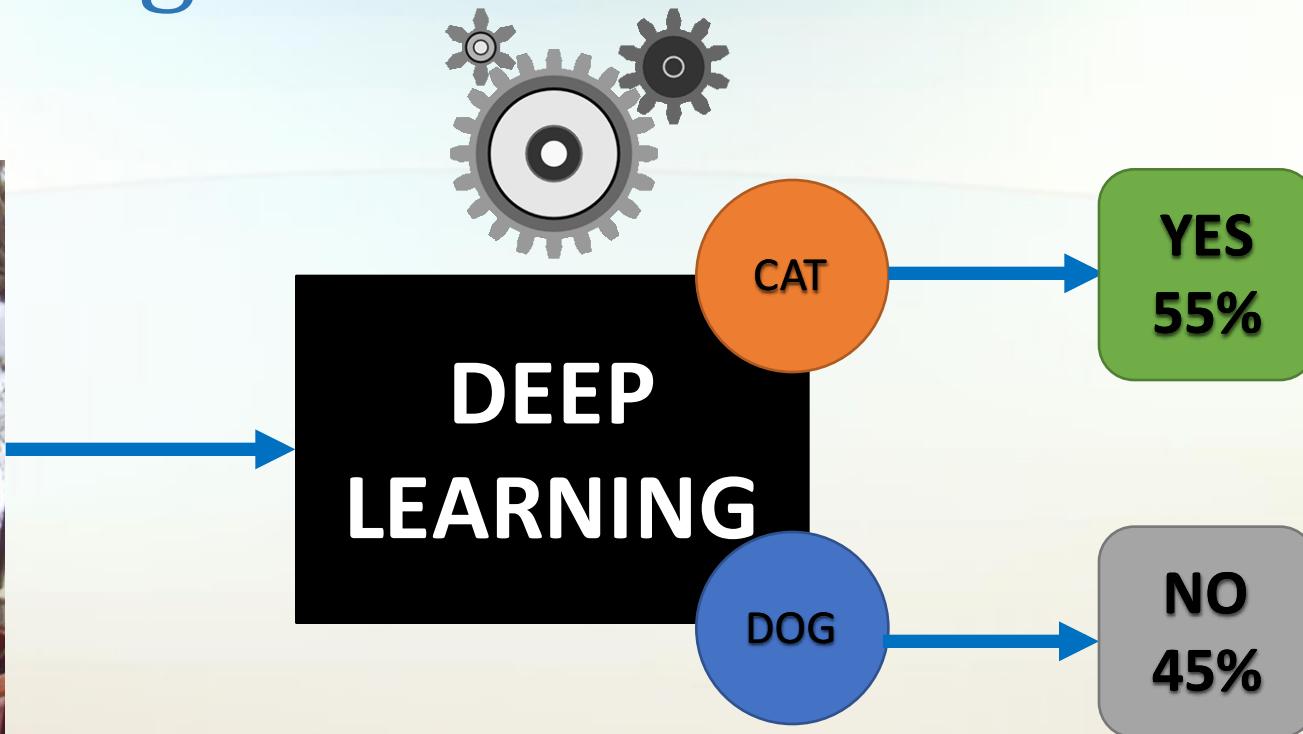
Cat/Dog Classifier



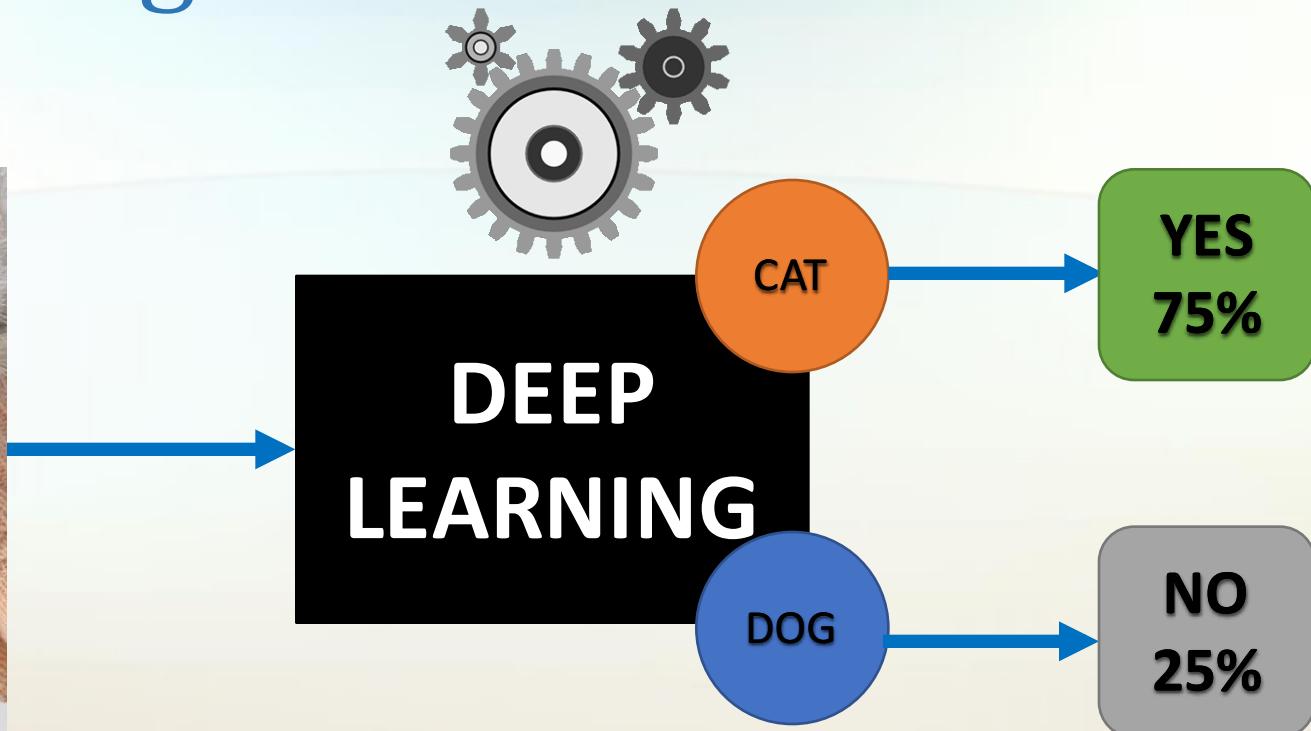
Cat/Dog Classifier



Cat/Dog Classifier



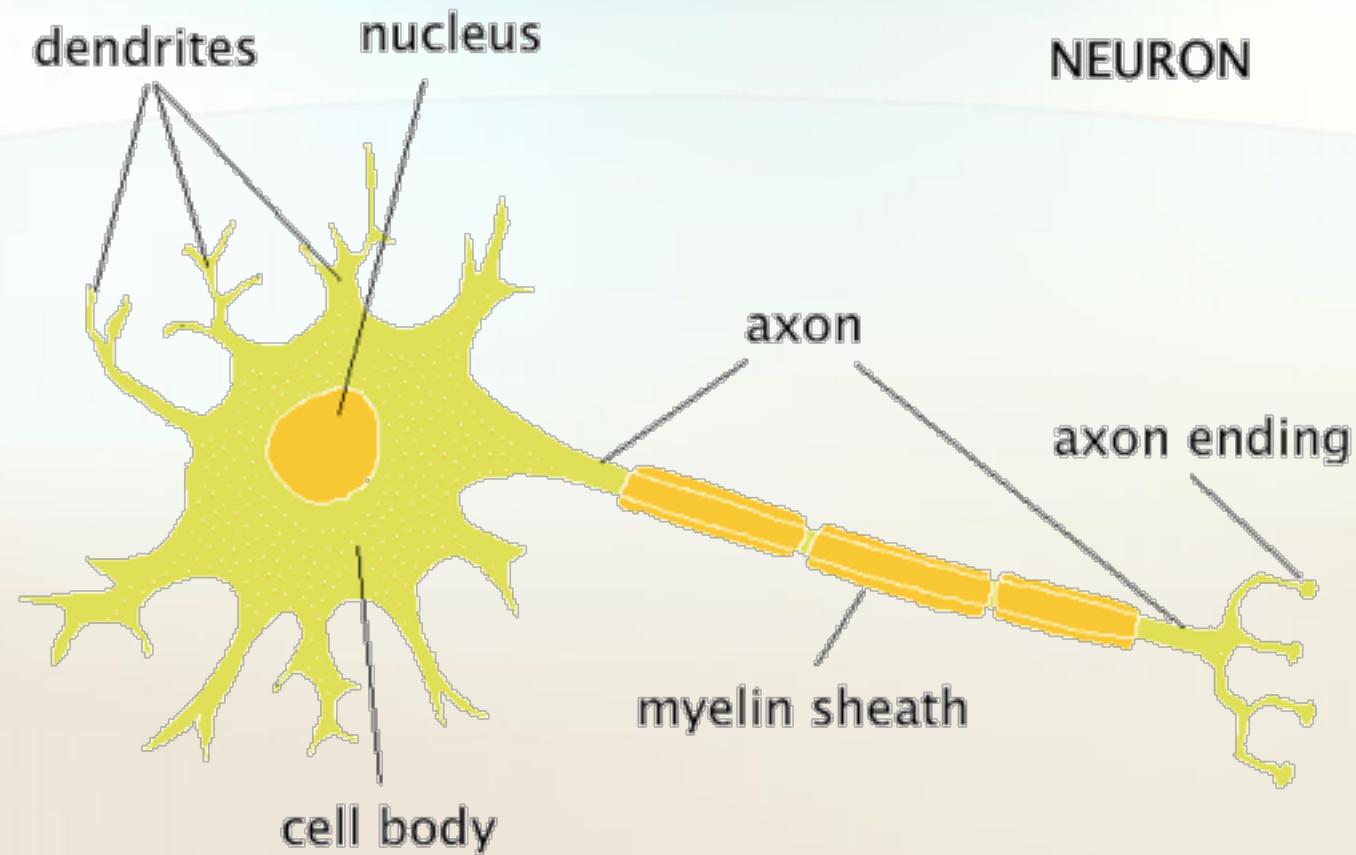
Cat/Dog Classifier

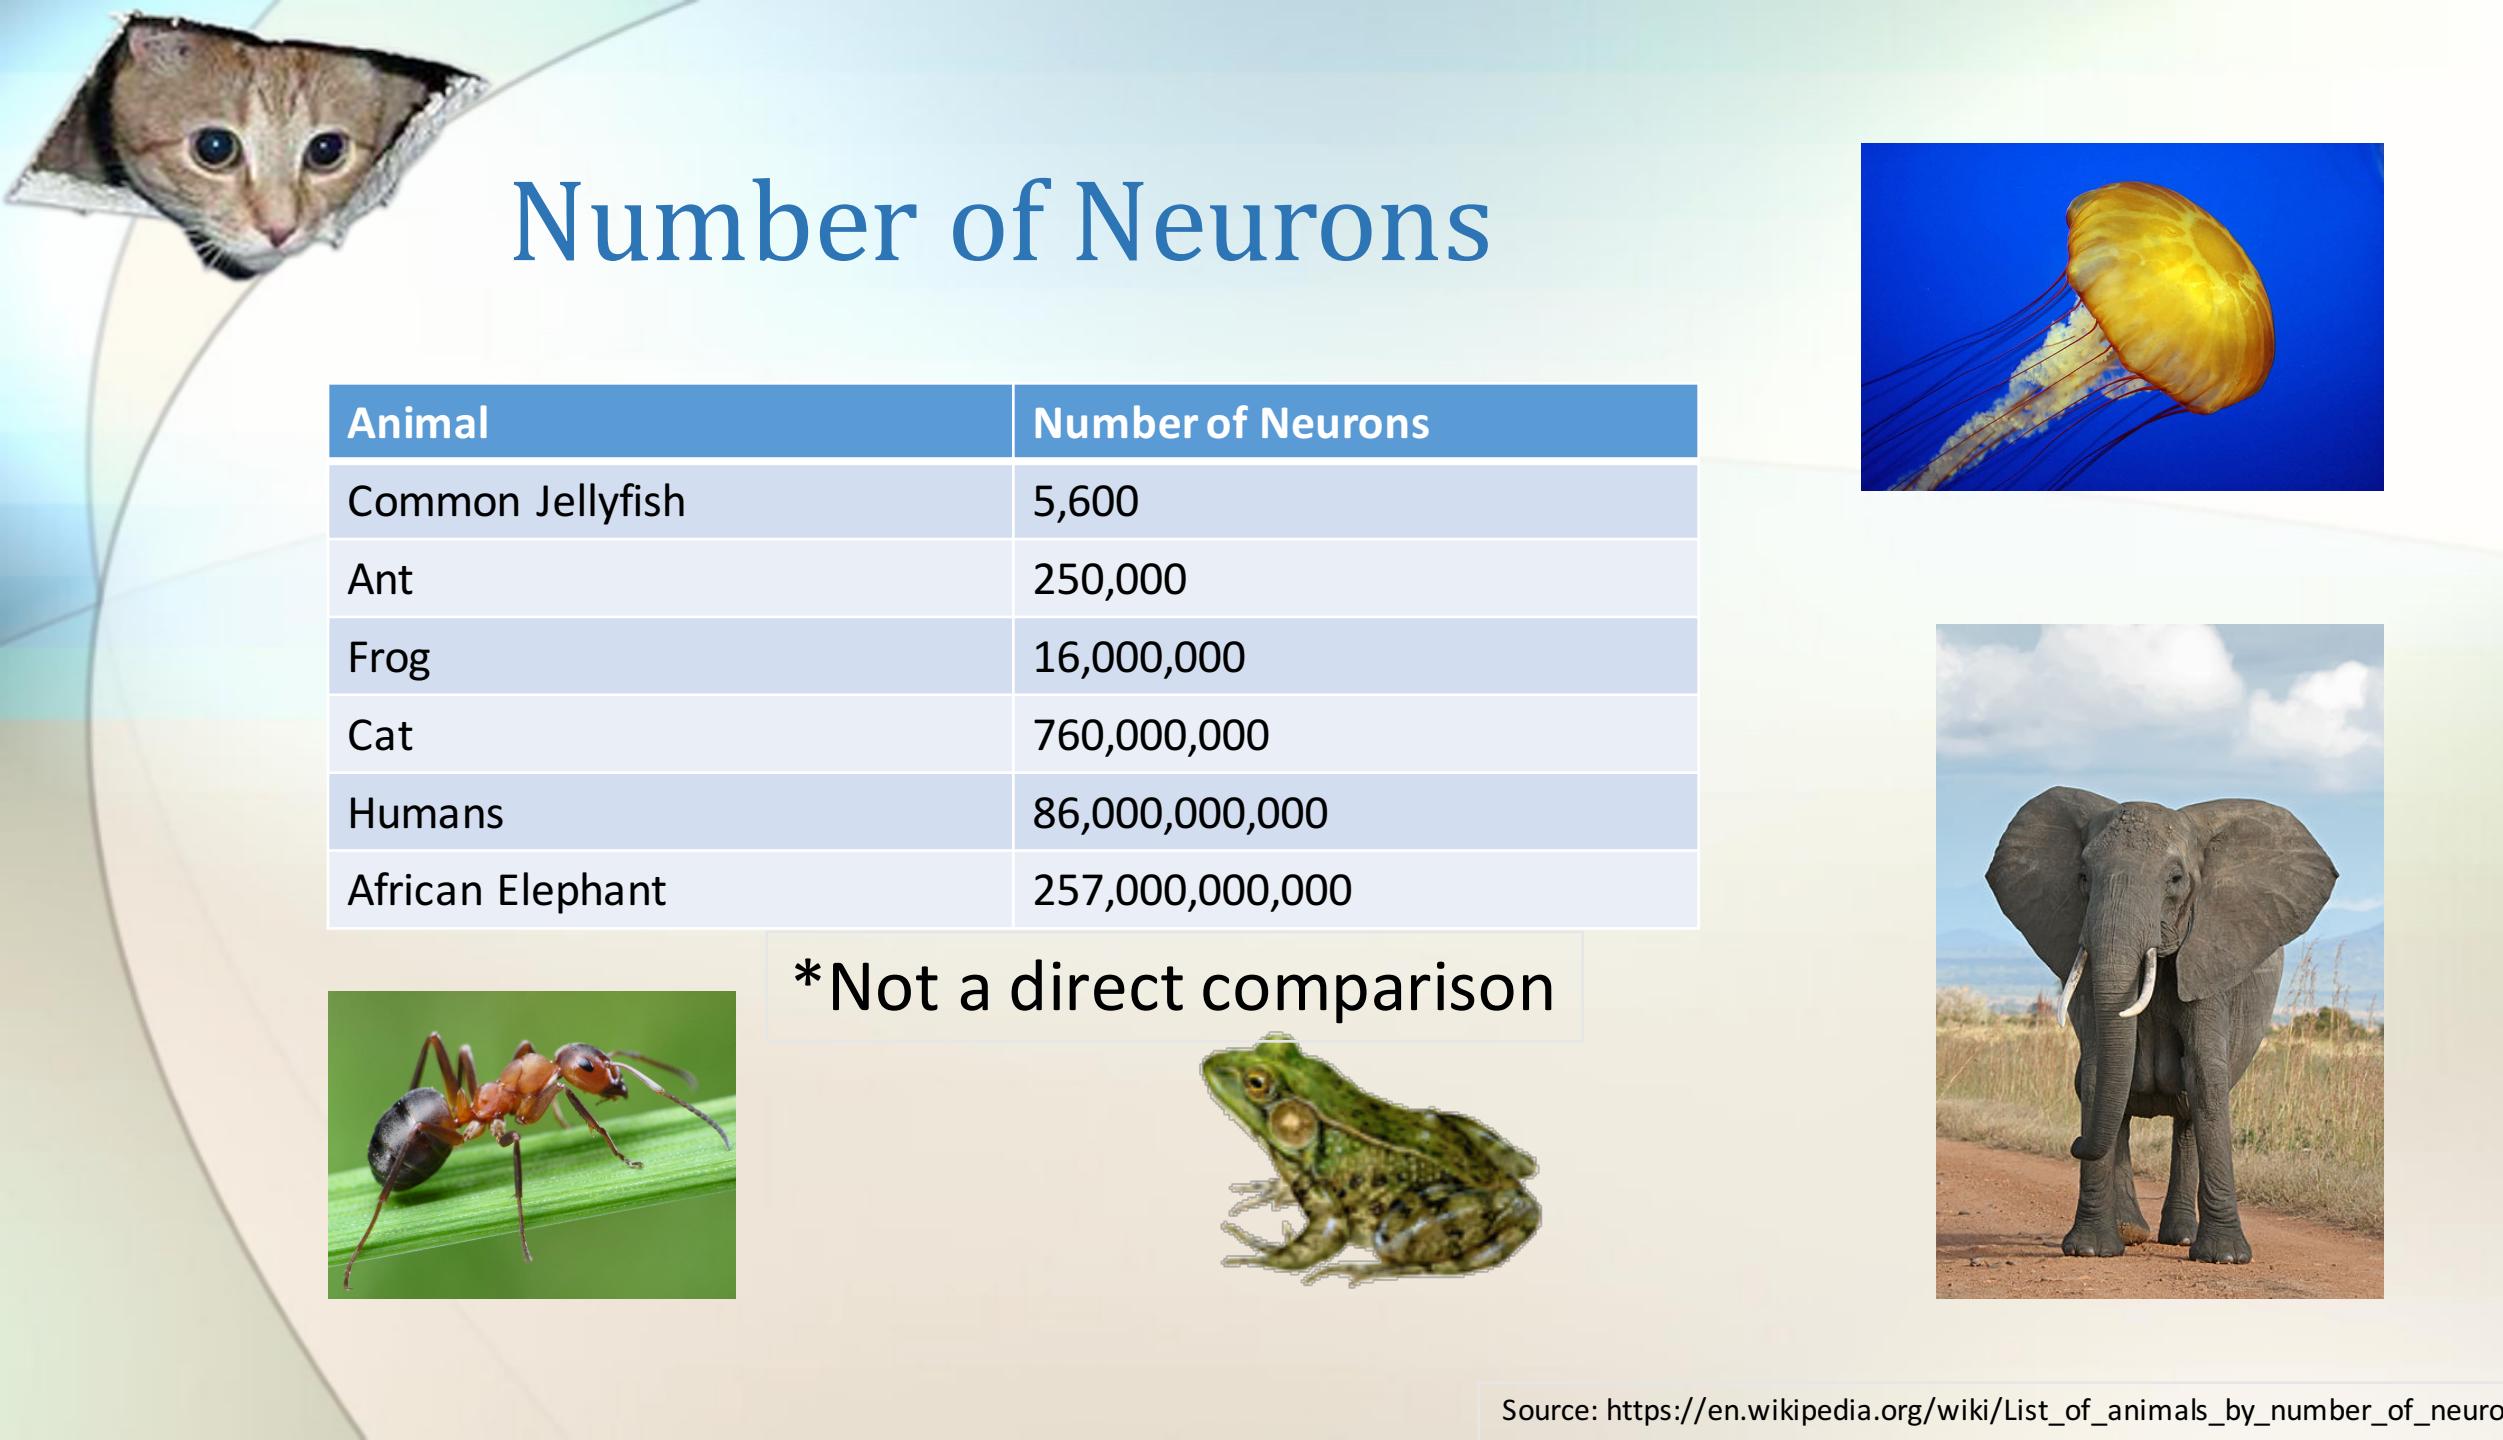


Topic Ordering (Logistics)

- There are many components in Artificial Neural Networks
- They all come together to make something useful.
- If something is not clear, please ask!
- Hard to figure out the best ordering of topics.

Neuron (Inspiration)





Number of Neurons

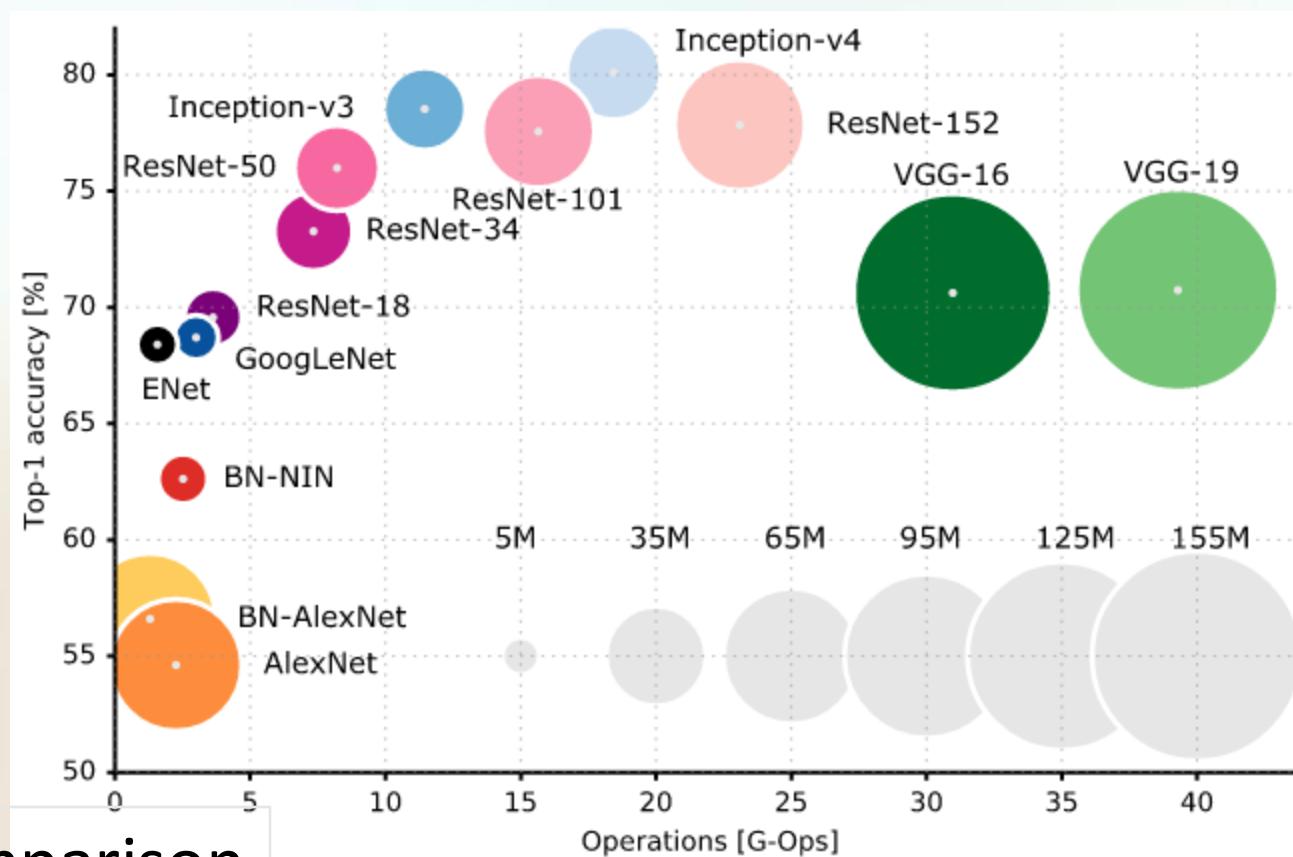
Animal	Number of Neurons
Common Jellyfish	5,600
Ant	250,000
Frog	16,000,000
Cat	760,000,000
Humans	86,000,000,000
African Elephant	257,000,000,000

*Not a direct comparison

Current Network Architectures

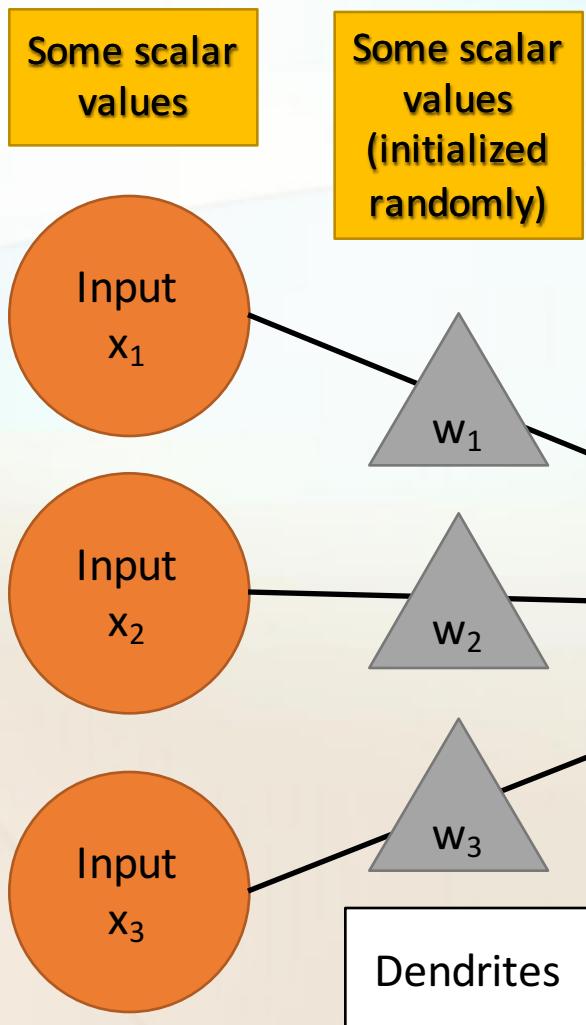
- Blob Size = Number of Parameters (connections between neurons)

ResNet-152 has
167,552 Neurons

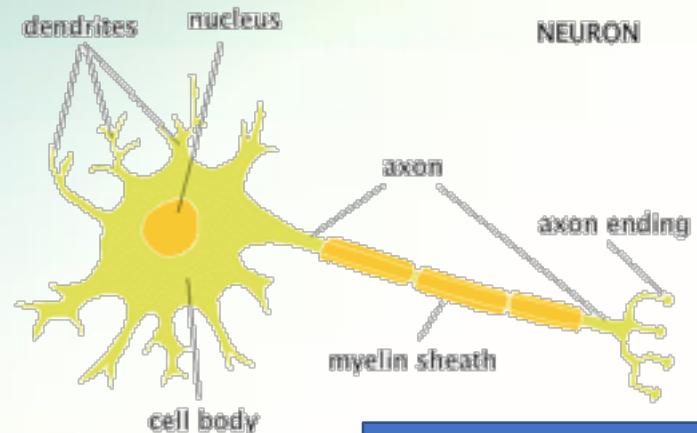


*Not a direct comparison

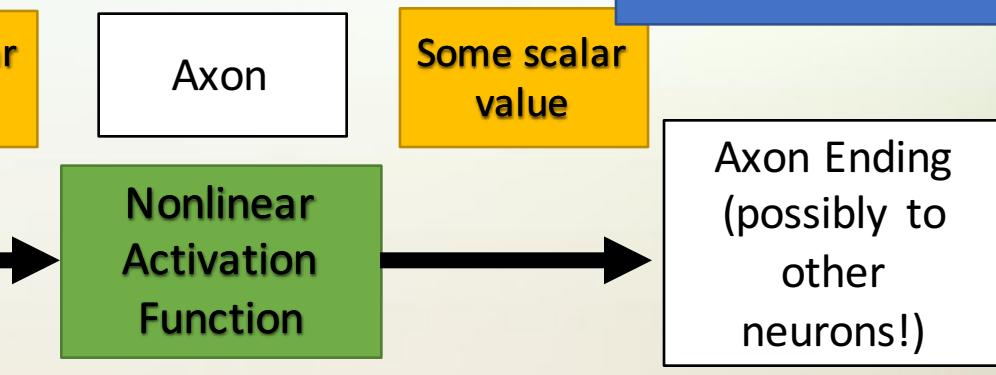
Artificial Neuron



$$\text{relu}([x_1 \ x_2 \ x_3] \begin{bmatrix} w_1 \\ w_2 \\ w_3 \end{bmatrix})$$



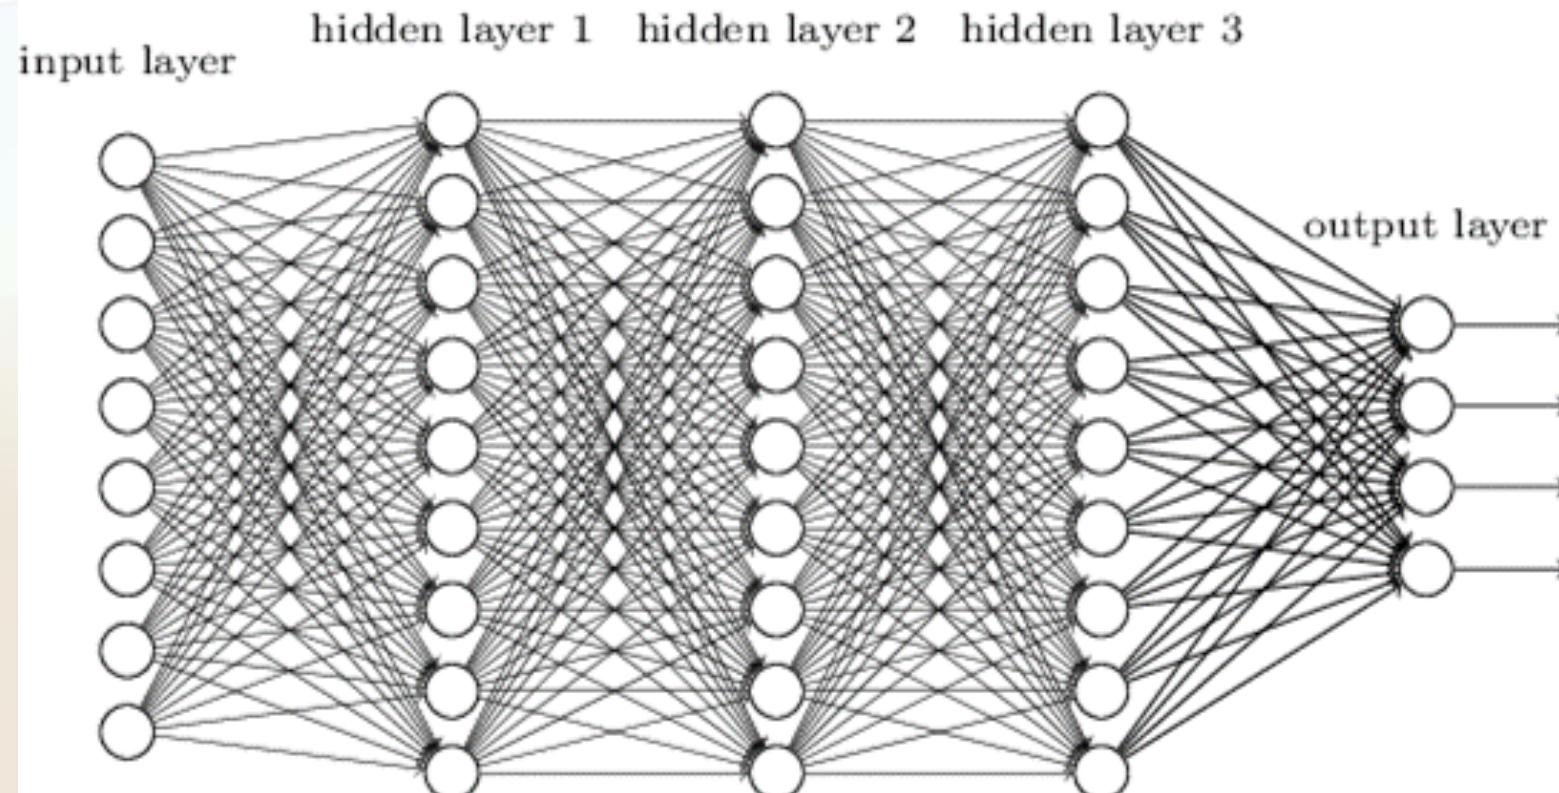
May use some output decision function



$$RELU(x) = \begin{cases} 0 & \text{if } x < 0 \\ x & \text{if } x \geq 0 \end{cases}$$

Combining Neurons Into Layers

Deep neural network



Backpropagation

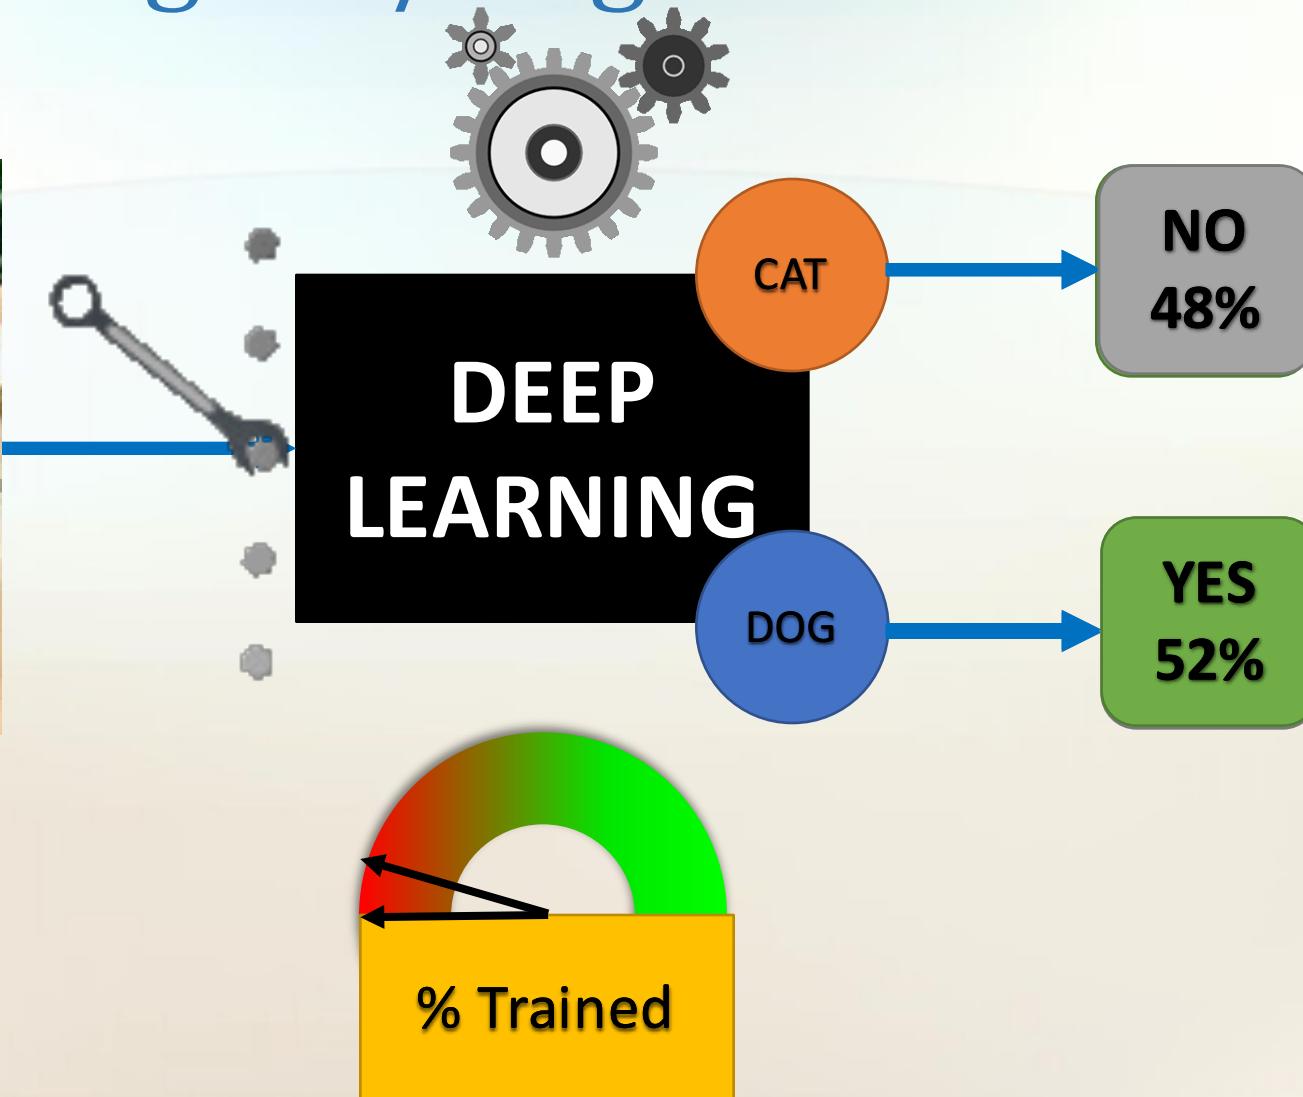
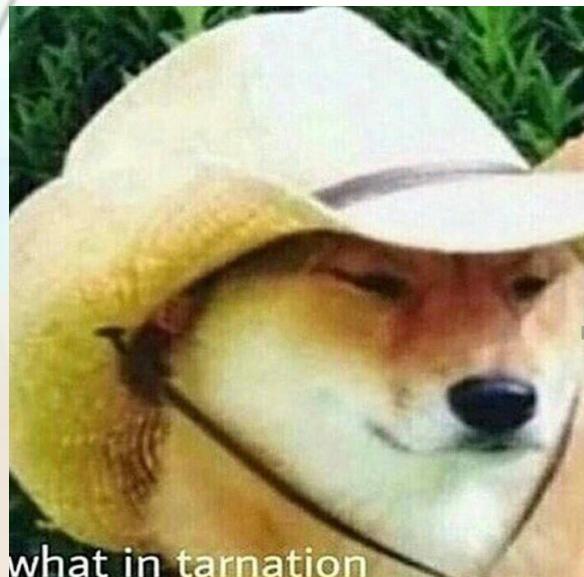
(How the network learns)

- How the network learns useful weights
- Will not go into depth on how it works
- **Don't** need to know it to use off-the-shelf components in TensorFlow
- **Do** need to know if you want to implement custom layers
- Basically through continuous differentiation (calculus), the network figures out how much each parameter (weights) contributed to an error and tweaks it to reduce the error.

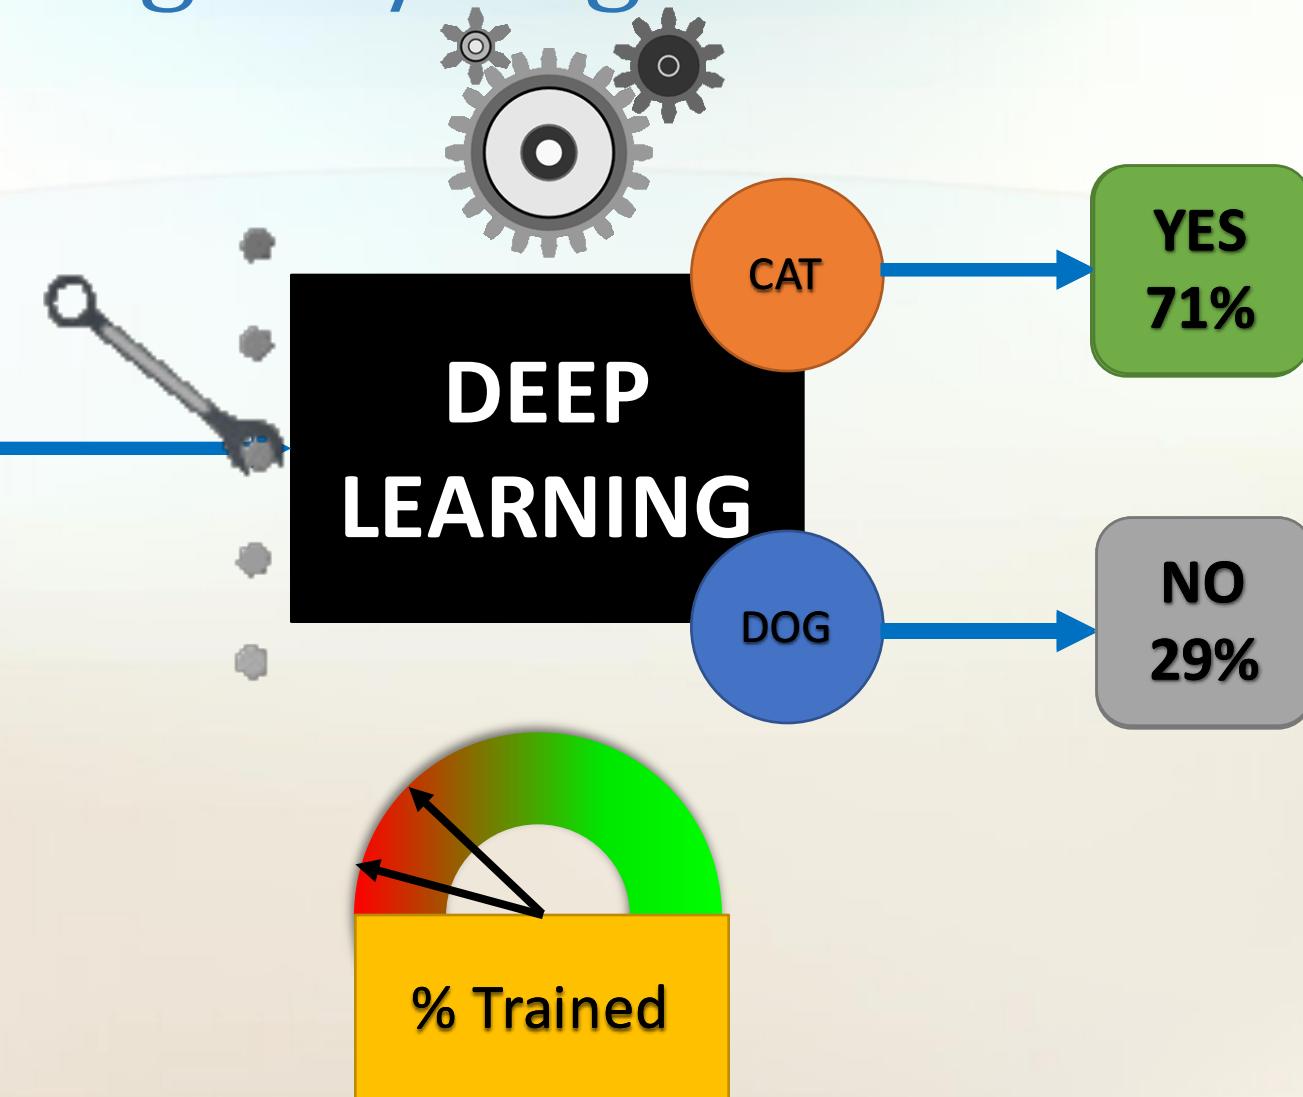
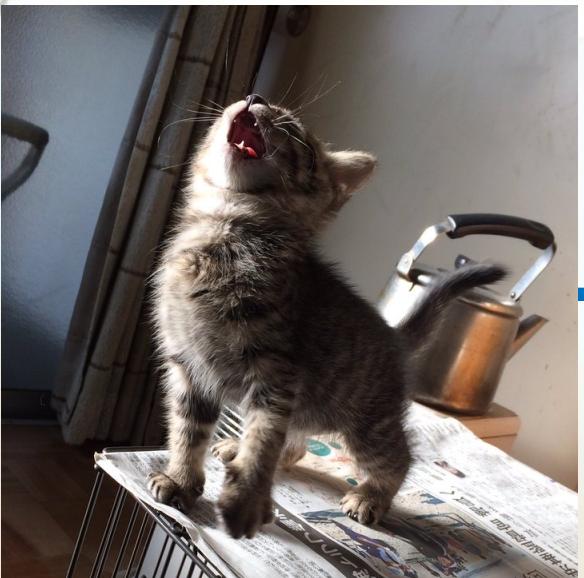
Training In A Nutshell

1. Forward pass some example input
2. Compare the network output with what the output **should** be
3. Backpropagation back through the network
4. Update weight values via Backpropagation
5. Now, whenever the network gets that input, the output should be closer to the goal output.

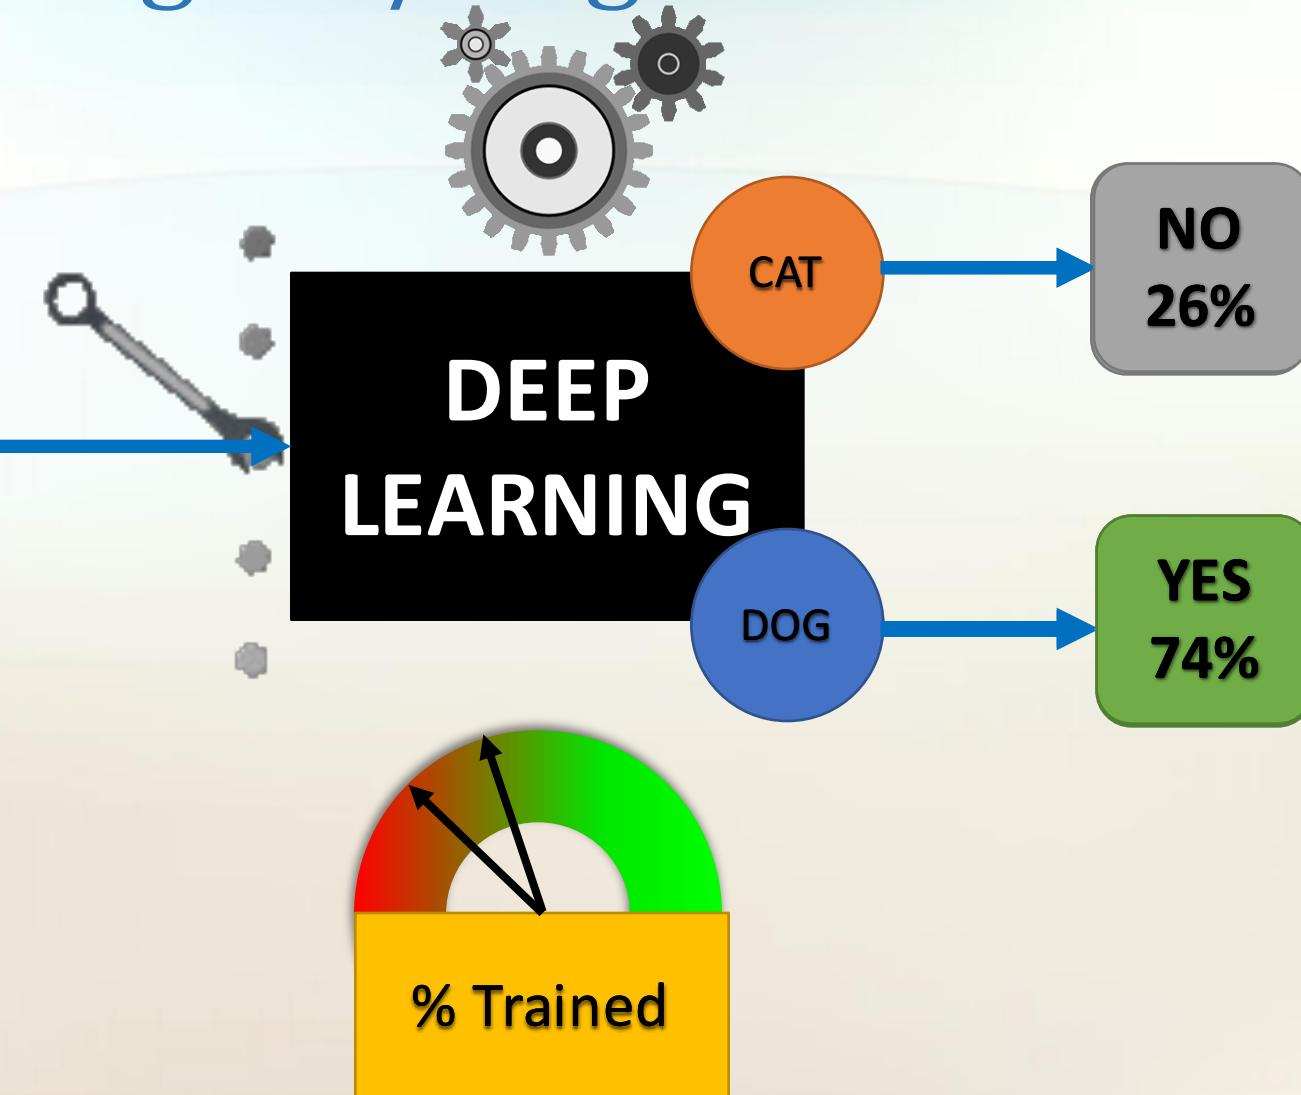
Training Cat/Dog Classifier



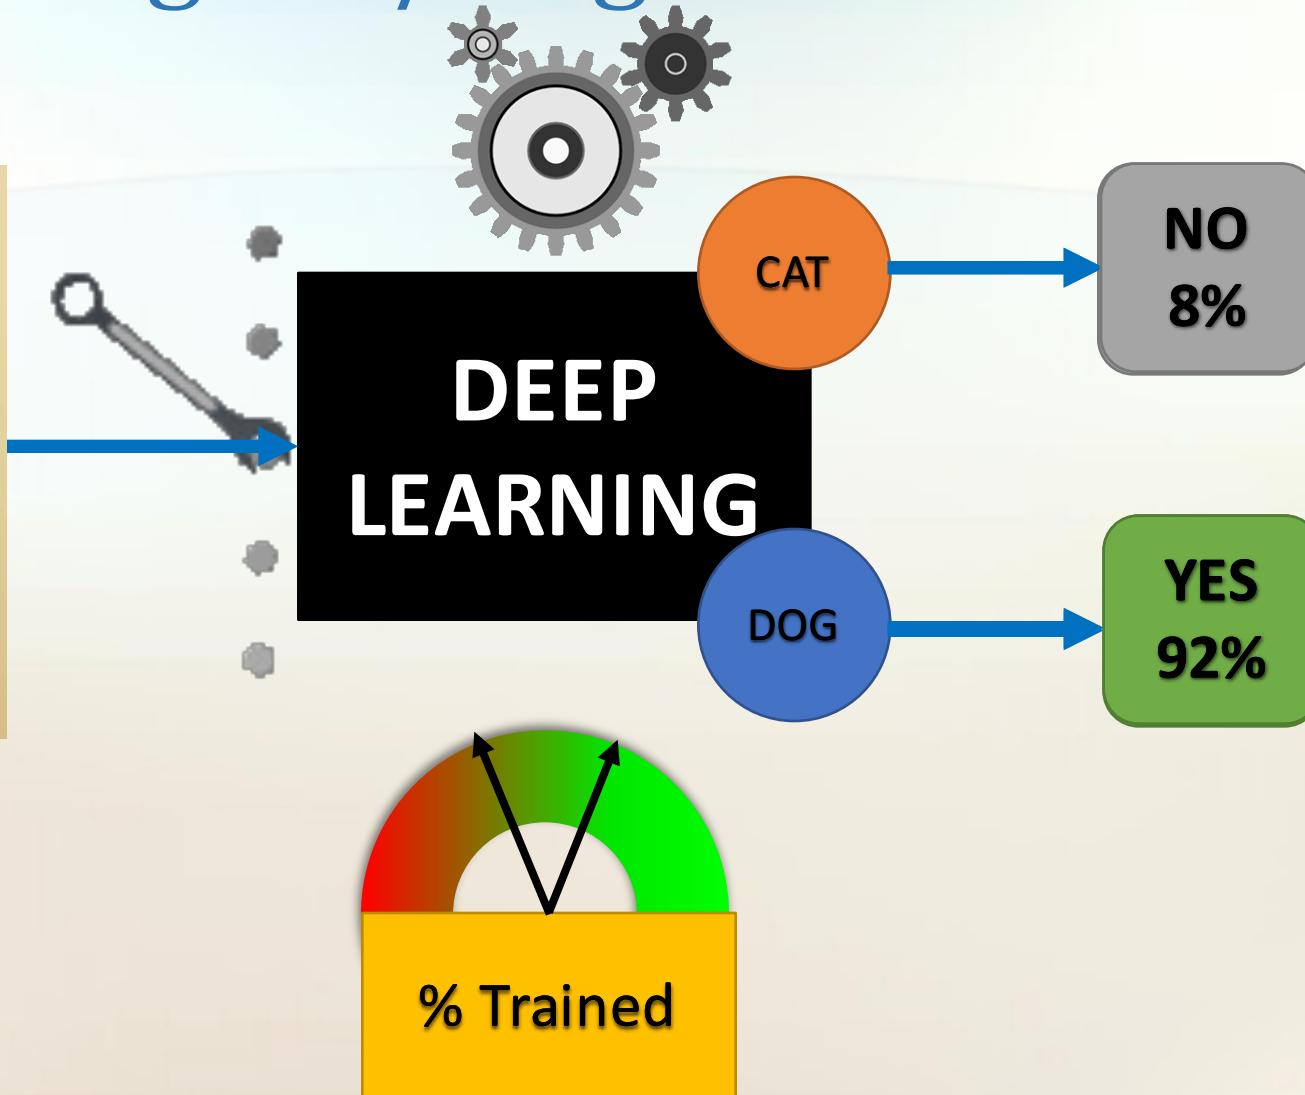
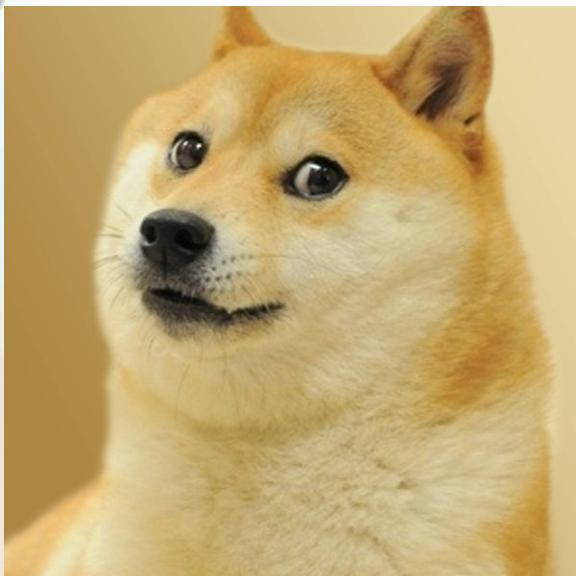
Training Cat/Dog Classifier



Training Cat/Dog Classifier



Training Cat/Dog Classifier



Training Cat/Dog Classifier

**DO THIS TENS/HUNDREDS
OF THOUSANDS OF TIMES**

GPUs! (Aside/Optional)



- Each neuron can be computed in parallel (independent)
- GPUs have hundreds (even thousands) of relatively weak cores
- Nvidia has a virtual monopoly (thanks to the CUDA Toolkit)
- Nvidia supplies AWS, Azure, etc.
- Google signed a deal with AMD (but also uses Nvidia)



NVIDIA.



AMD



TensorFlow!

- Google's publicly available Deep Learning library
- In competition with Caffe, Torch, etc
- Is becoming more and more popular in industry.
- Normally we leverage the power of GPUs (thousands of time speed up), but the installation for TensorFlow is Linux-only and more tedious than CPU-only installation.

Installation

- Download and Install [Anaconda](#) (Python 3.5)
- Download and install [TensorFlow](#) in a conda environment called “tensorflow”
- Download and install the Spyder IDE within the conda environment:

```
conda install -n tensorflow spyder
```

Importing the Library

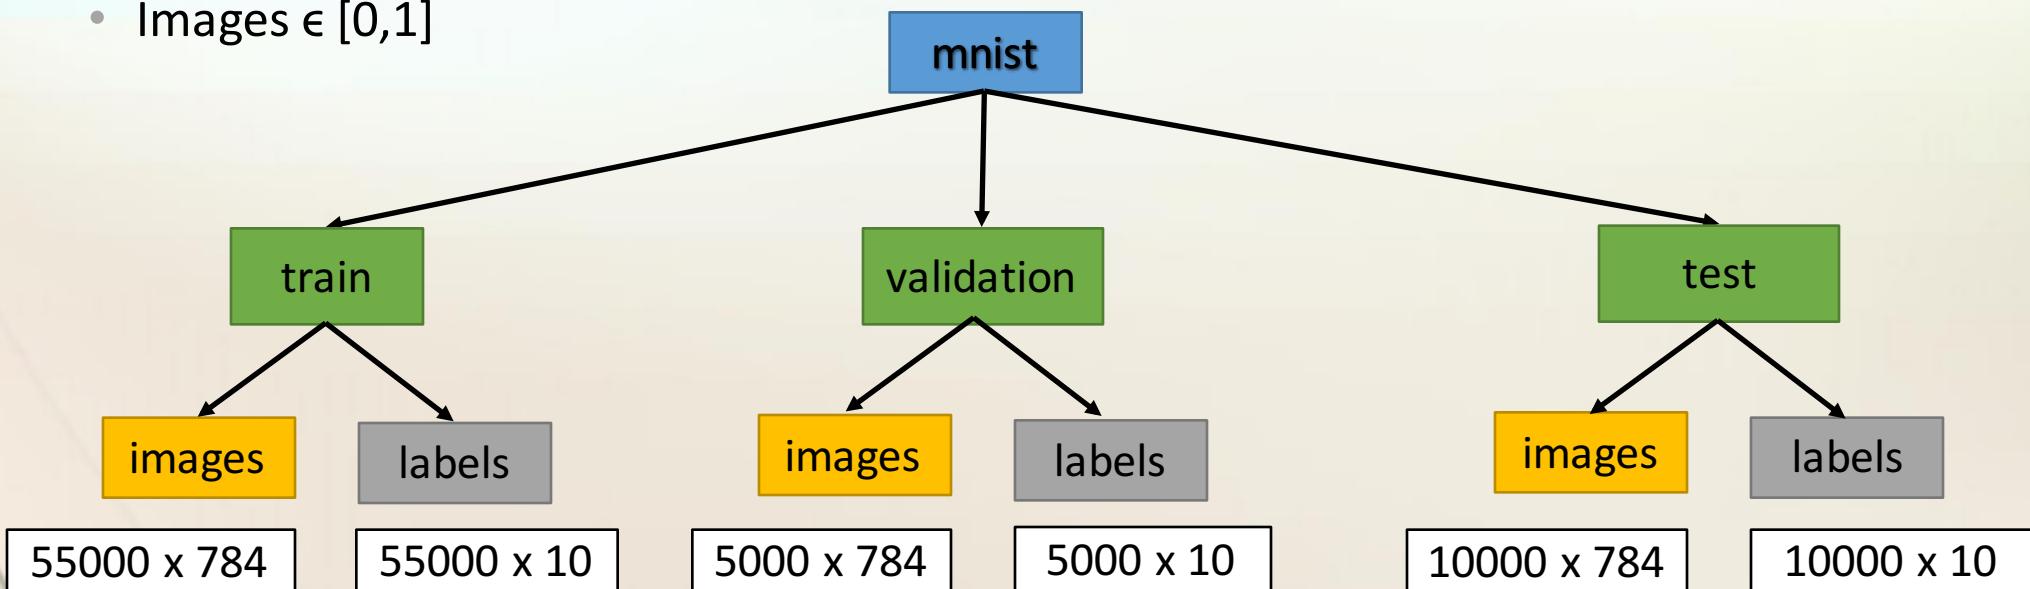
```
import tensorflow as tf
```

- Tensorflow functions can now be called using the tf prefix.
Example: tf.session()

Importing the MNIST Dataset

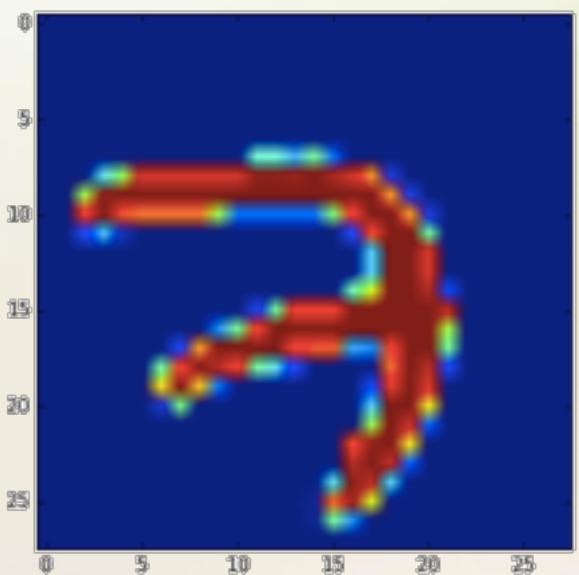
```
from tensorflow.examples.tutorials.mnist import input_data  
mnist = input_data.read_data_sets("MNIST_data/", one_hot=True)
```

- This downloads and loads in the MNIST digits dataset.
- Note: $784 = 28 \times 28$
- Images $\in [0,1]$



Viewing One MNIST Image

```
import matplotlib.pyplot as plt  
im = mnist.train.images[0, :]  
label = mnist.train.labels[0, :]  
im = im.reshape([28, 28])
```



Softmax and One-Hot Encoding

- We want the network to output a percent certainty it believes some image belongs to some label.
- Softmax remaps the output layers to percentages.

$$\text{softmax}(\nu_j) = \frac{e^{\nu_j}}{\sum_i e^{\nu_i}}$$

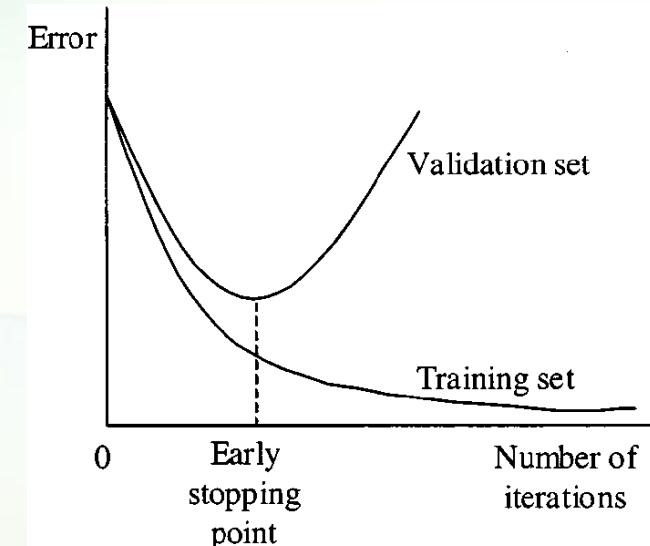
- Example: [1.0, 1.0, 6.0] -> [0.0066, 0.0066, 0.9867]
- One-Hot Encoding (with 10 options): 3 -> [0 0 0 1 0 0 0 0 0 0]

Train, Validation, Test

- Typical Data Breakup:

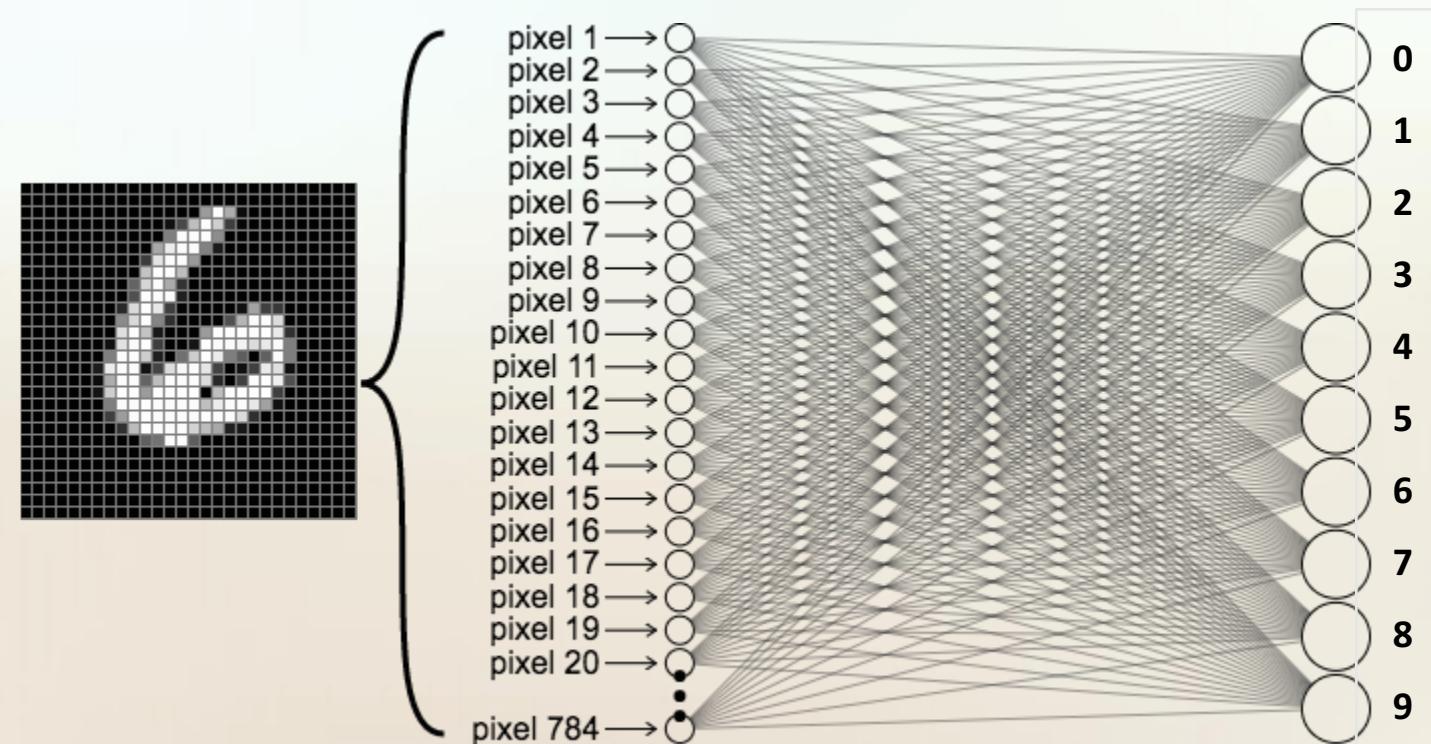
	Train	Validation	Test
Percentage:	80~90	5~10	7~15

- Important that every dataset is representative
- Train on the training dataset.
- Check performance during training with Validation dataset
- See actual network performance with Test dataset.



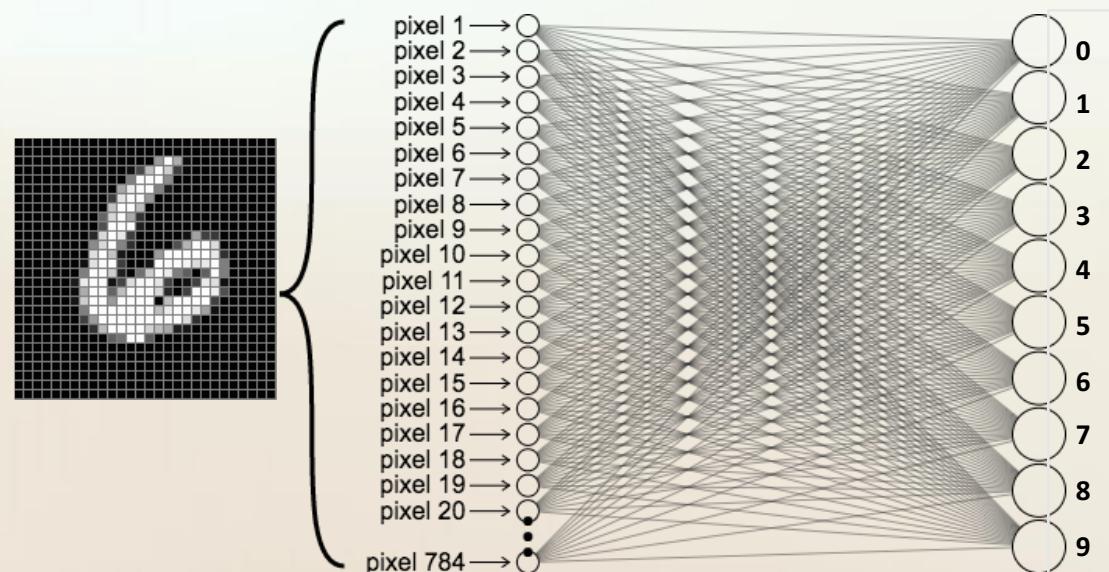
Define Model (Not So Deep Learning)

- We will construct a feedforward fully connected neural network that is 1 layer deep.

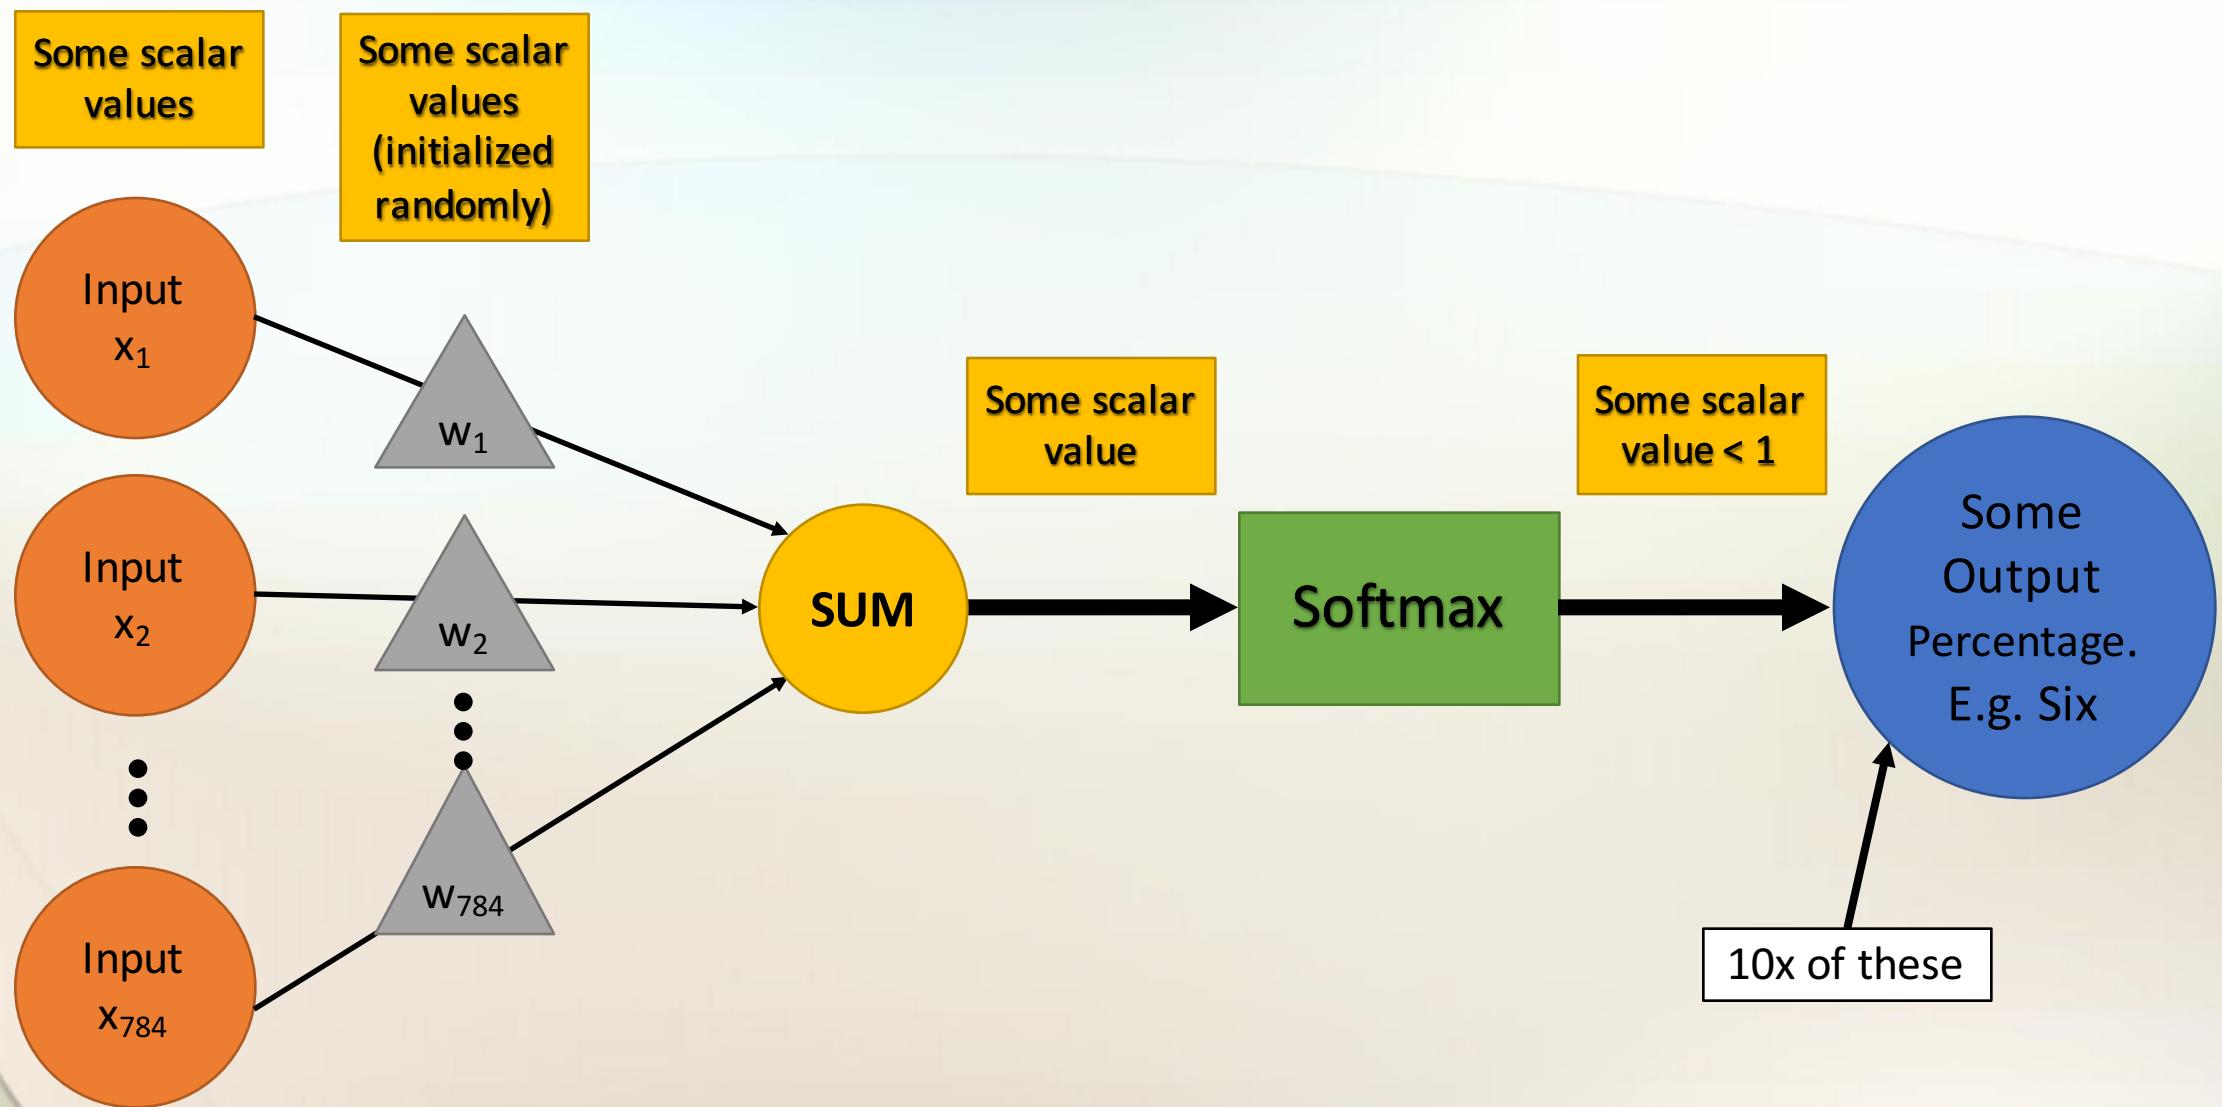


Define Our Model

- Each pixel will have 10 associated weights (1 for each layer)
- There are 784 pixels in each image
- Our weight matrix will have dimensions 784 by 10



Artificial Neuron (Simplified)

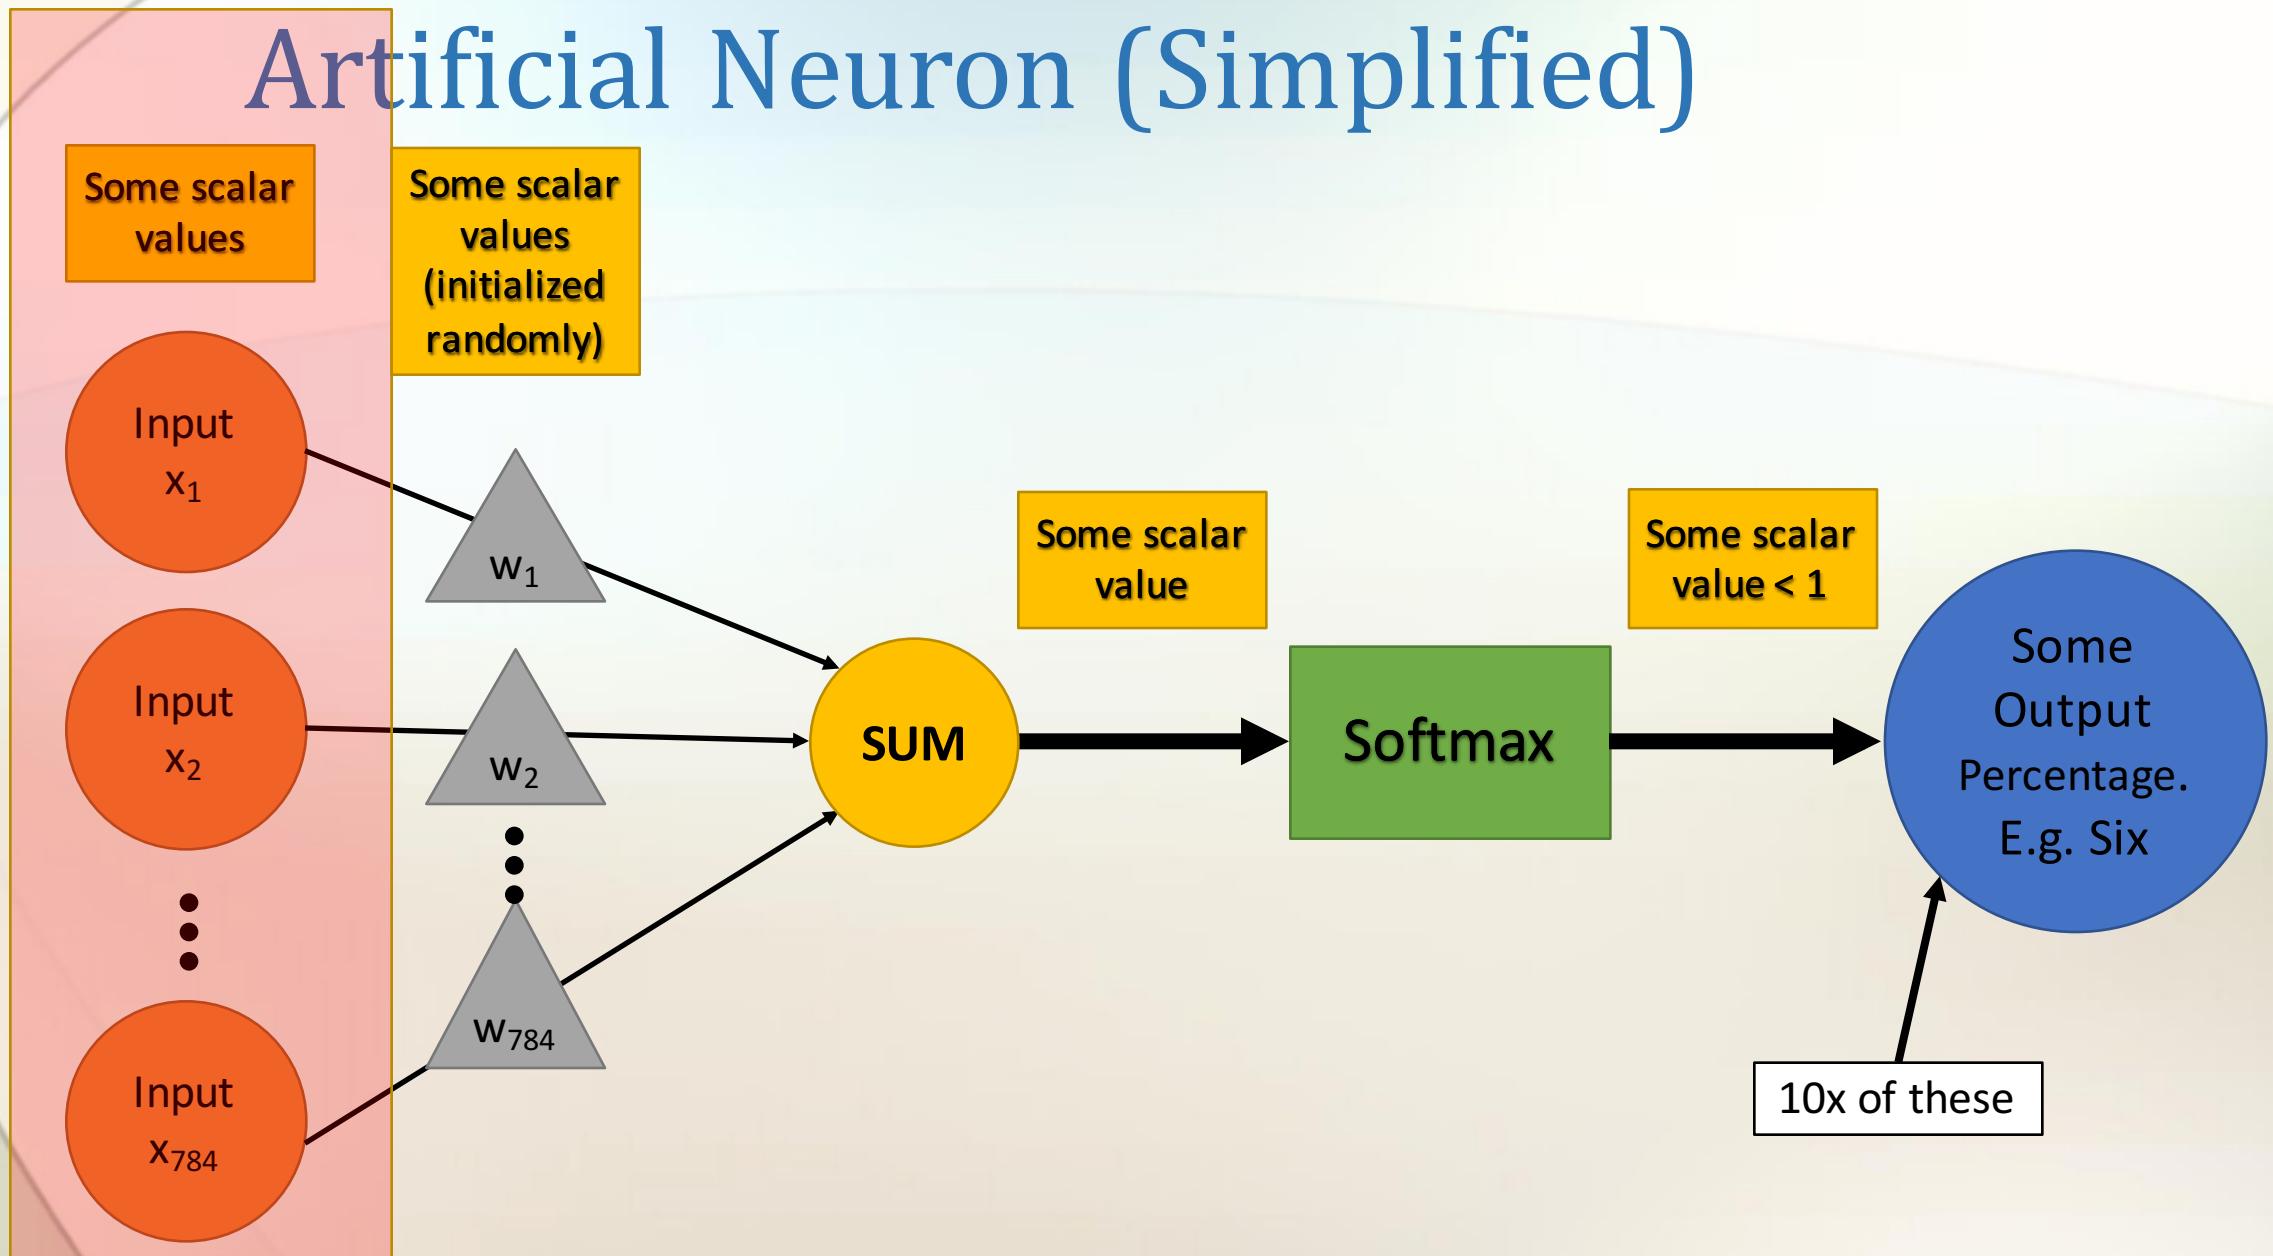


Input Variable

```
x = tf.placeholder(tf.float32, [None, 784])
```

- Creates a placeholder variable “x”
 - “x” doesn’t have a specific value yet
 - It’s just a variable, like in math
- Placeholder for our input images
- It is of type “TensorFlow Float 32”
- It has shape “None” by 784
 - None means the first dimension can have any length
 - 784 is the size of one image

Artificial Neuron (Simplified)

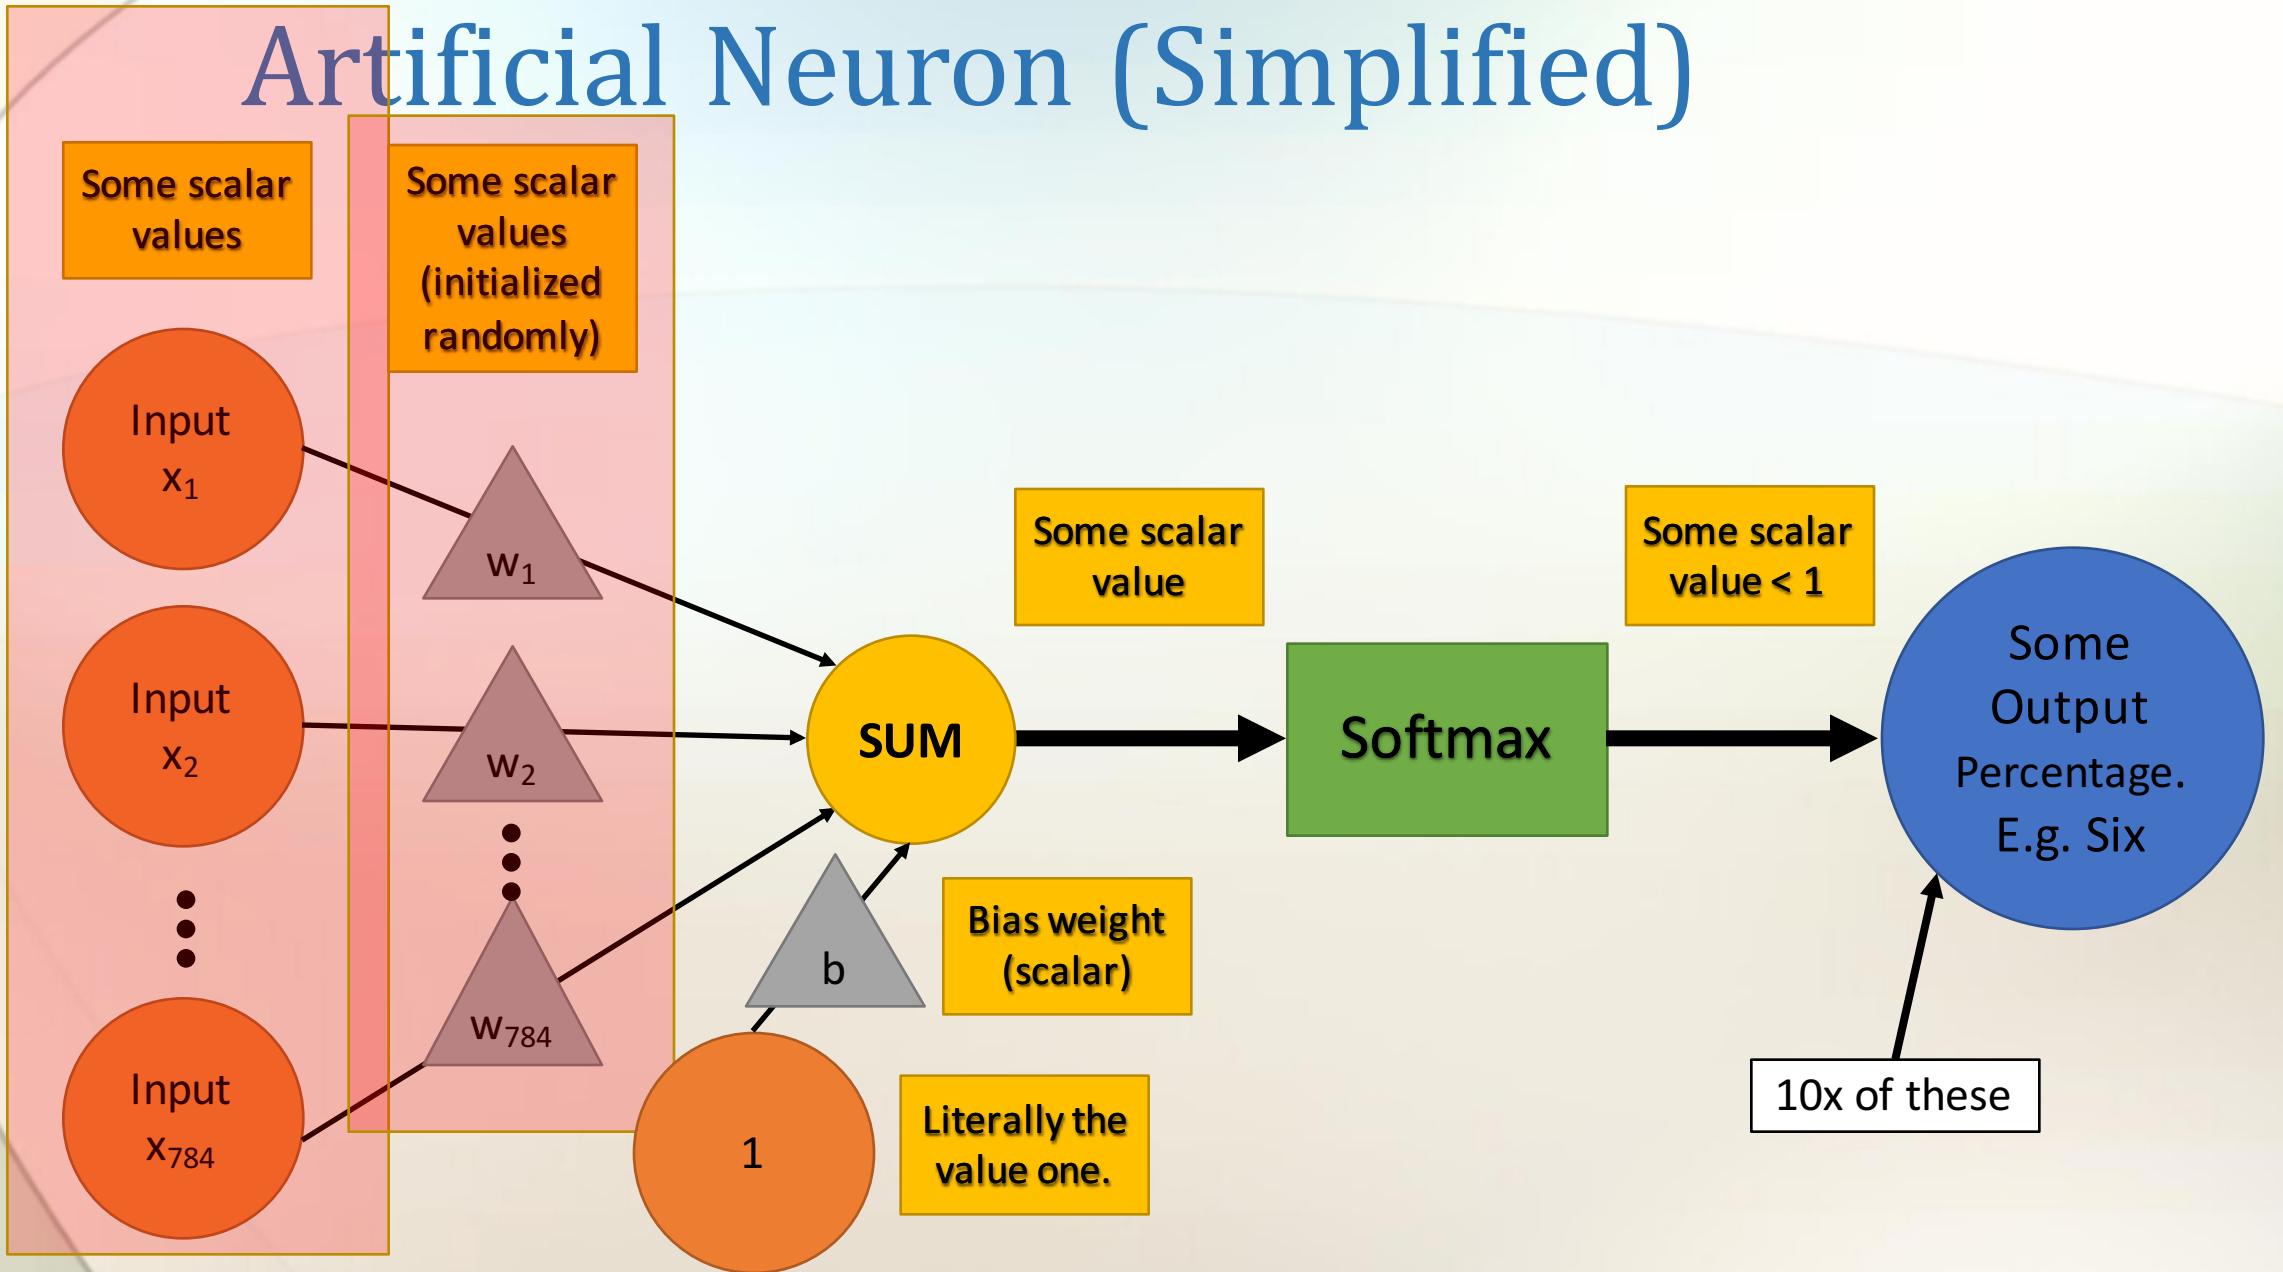


Network Variables (Weights)

```
W = tf.Variable(tf.zeros([784, 10]))
```

- Creates a variable W (for “weight”) of size 784 by 10
 - All elements of W are set to 0
 - Unlike “placeholder”, Variable contains determined values

Artificial Neuron (Simplified)

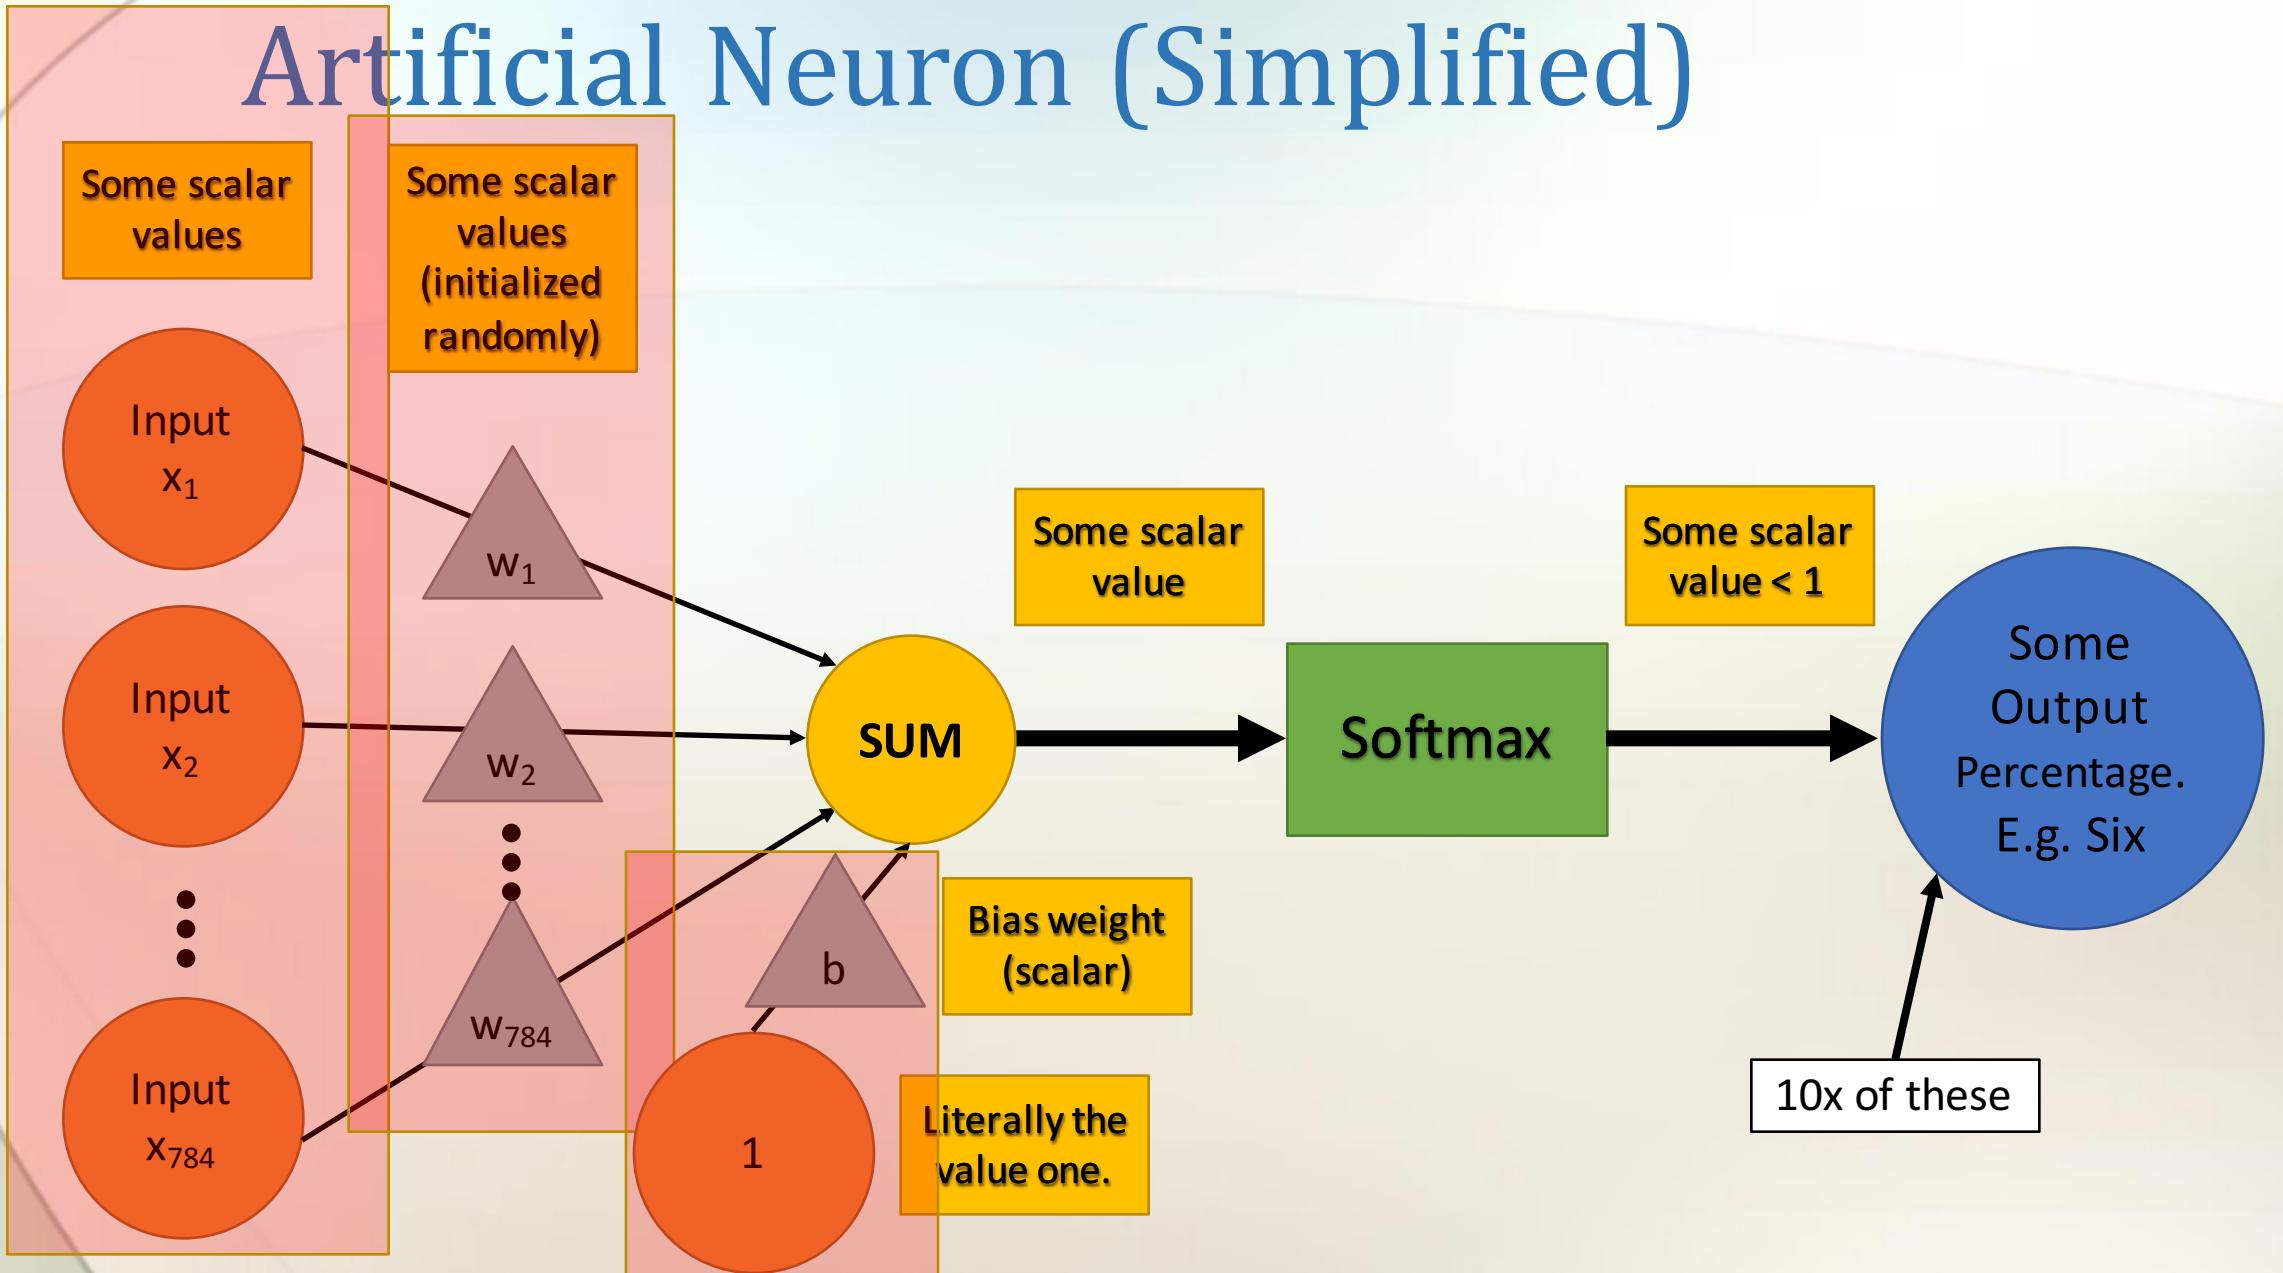


Network Variables (Biases)

```
b = tf.Variable(tf.zeros([10]))
```

- Creates a variable b (for “bias”) of size 10 (by 1)
 - All elements of b are set to 0
 - Unlike “placeholder”, Variable contains determined values

Artificial Neuron (Simplified)

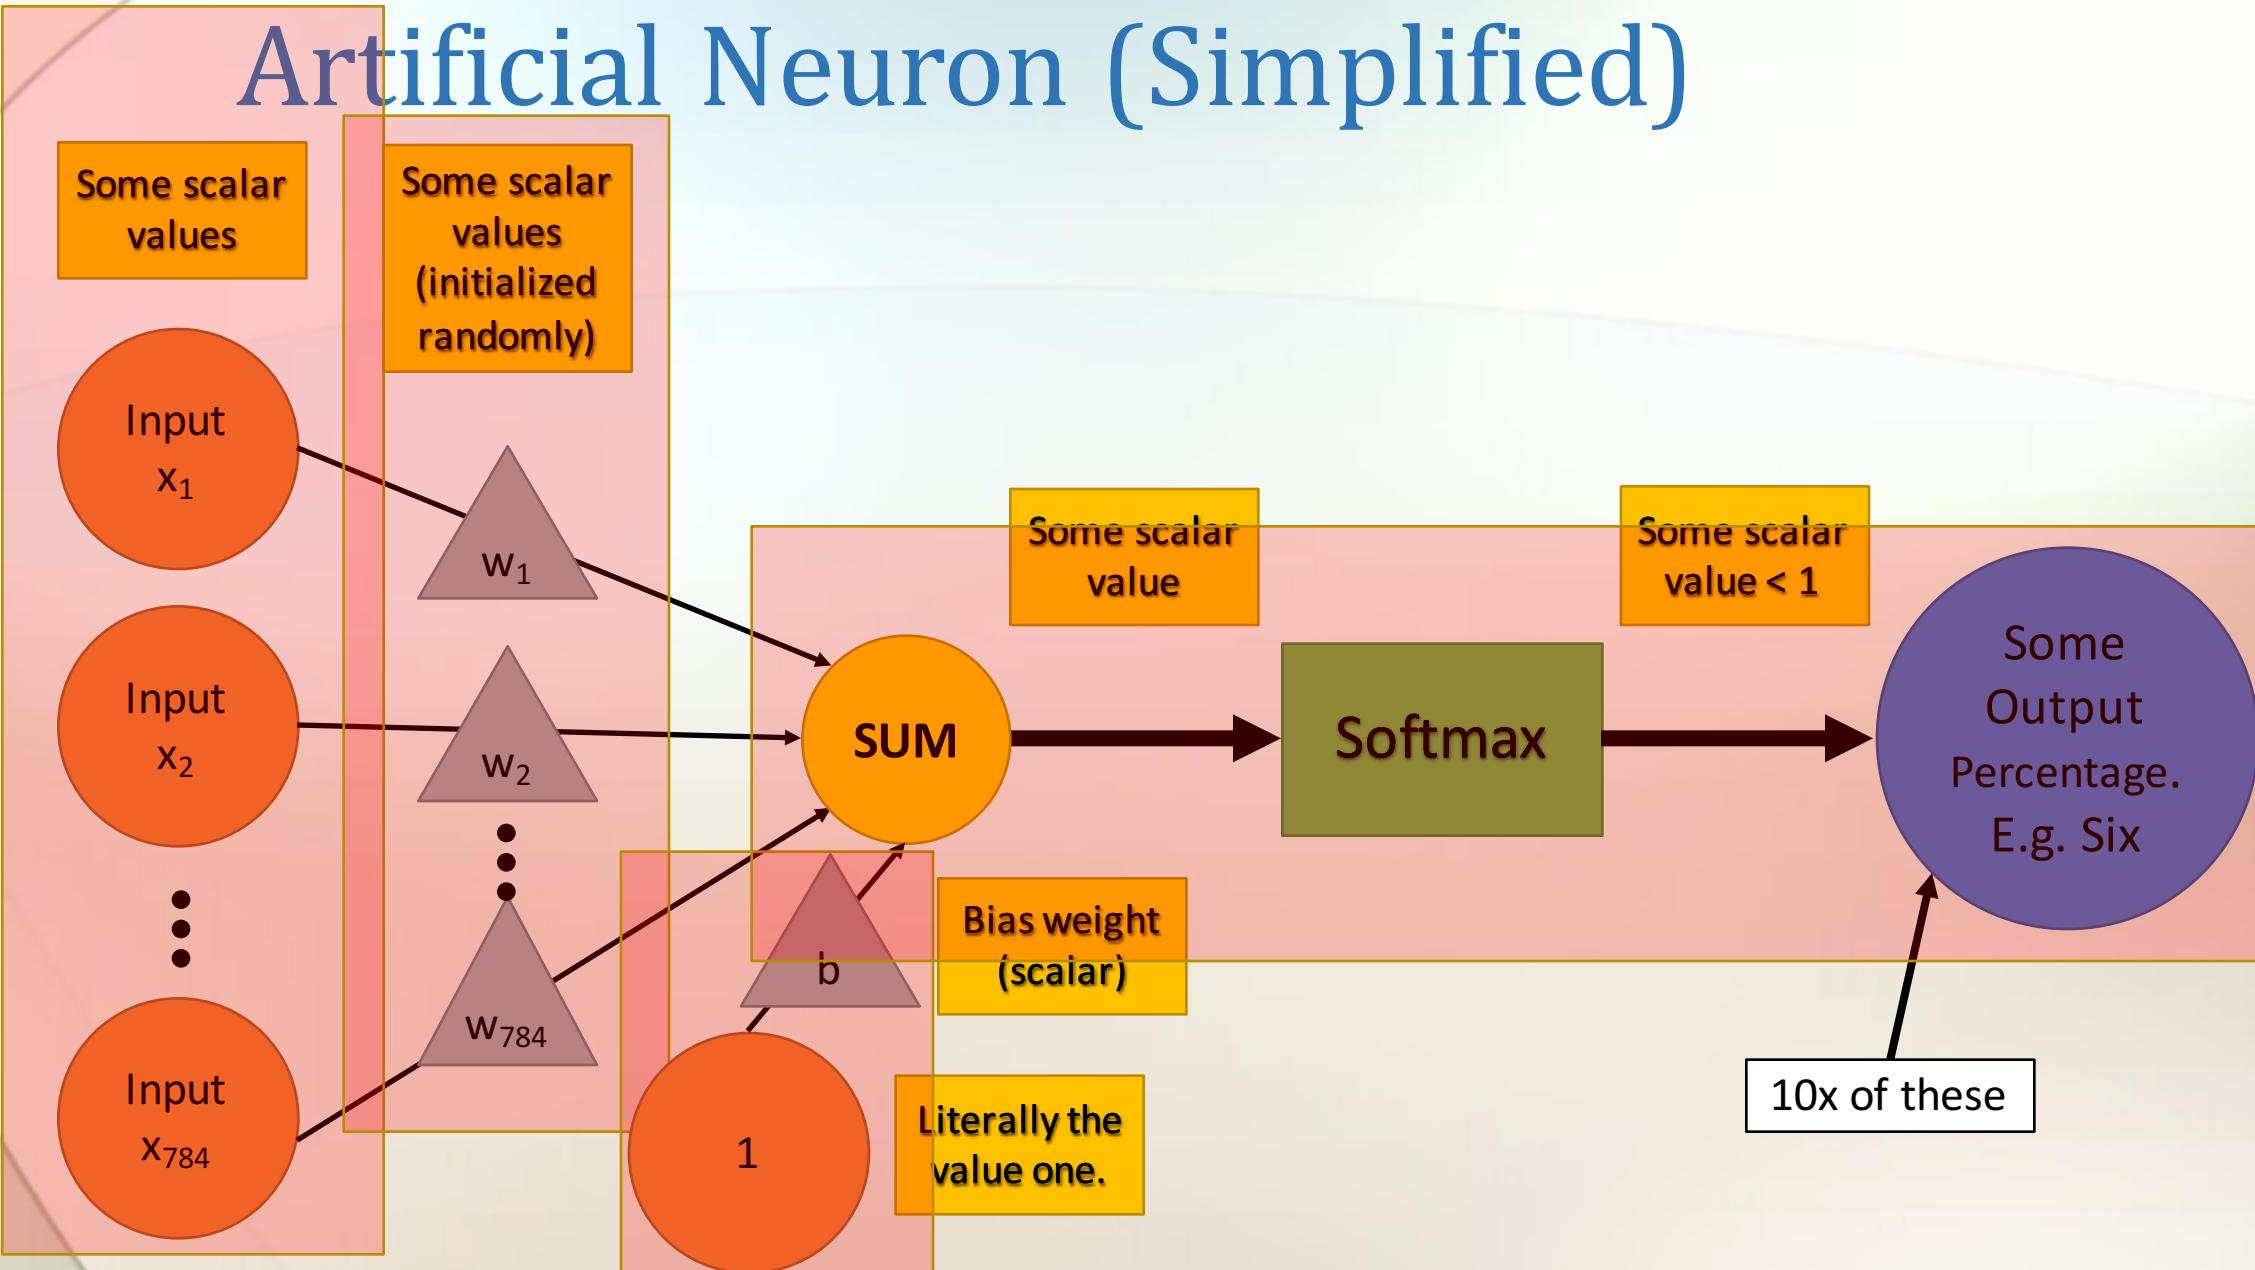


Network Output Variables

```
y = tf.nn.softmax(tf.matmul(x, W) + b)
```

- `tf.matmul(x, W)` performs a matrix multiplication between input variable “x” and weight variable W
- `tf.matmul(x, W) + b` add the bias variable
- `tf.nn.softmax(tf.matmul(x, W) + b)` perform the softmax operation
- `y` will have dimension None by 10

Artificial Neuron (Simplified)



Ground Truth Output Variables

```
yTruth = tf.placeholder(tf.float32, [None, 10])
```

- Creates a placeholder variable “yTruth”
 - “y” doesn’t have a specific value yet
 - It’s just a variable, like in math
- Placeholder for Ground Truth one-hot label outputs
- It is of type “TensorFlow Float 32”
- It has shape “None” by 10
 - None means the first dimension can have any length
 - 10 is the number of classes

Loss Variable

```
loss = tf.reduce_mean(-tf.reduce_sum(yTruth * tf.log(y),  
                                reduction_indices=1))
```

- `tf.log(y)` turns values close to 1 to be close to 0, and values close to 0 to be close to –infinity
- `yTruth*tf.log(y)` only keeps the value of the actual class
- `-tf.reduce_sum(yTruth*tf.log(y), reduction_indices=1)`
Sums along the class dimension (mostly 0's), fixes the sign
- `tf.reduce_mean(...)` averages the vector into a scalar

Loss Variable Example

Predict (y)		Average	Ground Truth (yTruth)	
Cat	Dog	Loss	Cat	Dog
0.25	0.75		0	1
0.90	0.10	0.4363	1	0
0.6	0.4		0	1

Implement Training Algorithm

```
lr = 0.5 # learning rate  
trainStep = tf.train.GradientDescentOptimizer(lr).minimize(loss)
```

- Learning rate is how much to proportionally change weights per training example.
- Minimize the loss function
- *Magic*

Begin the TensorFlow Session

- Up to this point, we have just been laying down a blueprint for TensorFlow to follow, but it hasn't "built" anything yet.

```
init = tf.global_variables_initializer()  
sess = tf.Session()  
sess.run(init)
```

- Initialize variables
- Create and run a TensorFlow session.

Run Through the Training Dataset

```
batchSize = 100
for i in range(1000):
    # get some images and their labels
    xBatches, yBatches = mnist.train.next_batch(batchSize)
    sess.run(trainStep, feed_dict={x:xBatches, yTruth:yBatches})
```

- Repeat 1000 times
- Gets a small random (100) subset of our training dataset
- Train on that small subset (this line updates the weights)
- Hopefully have a trained network once its done looping!

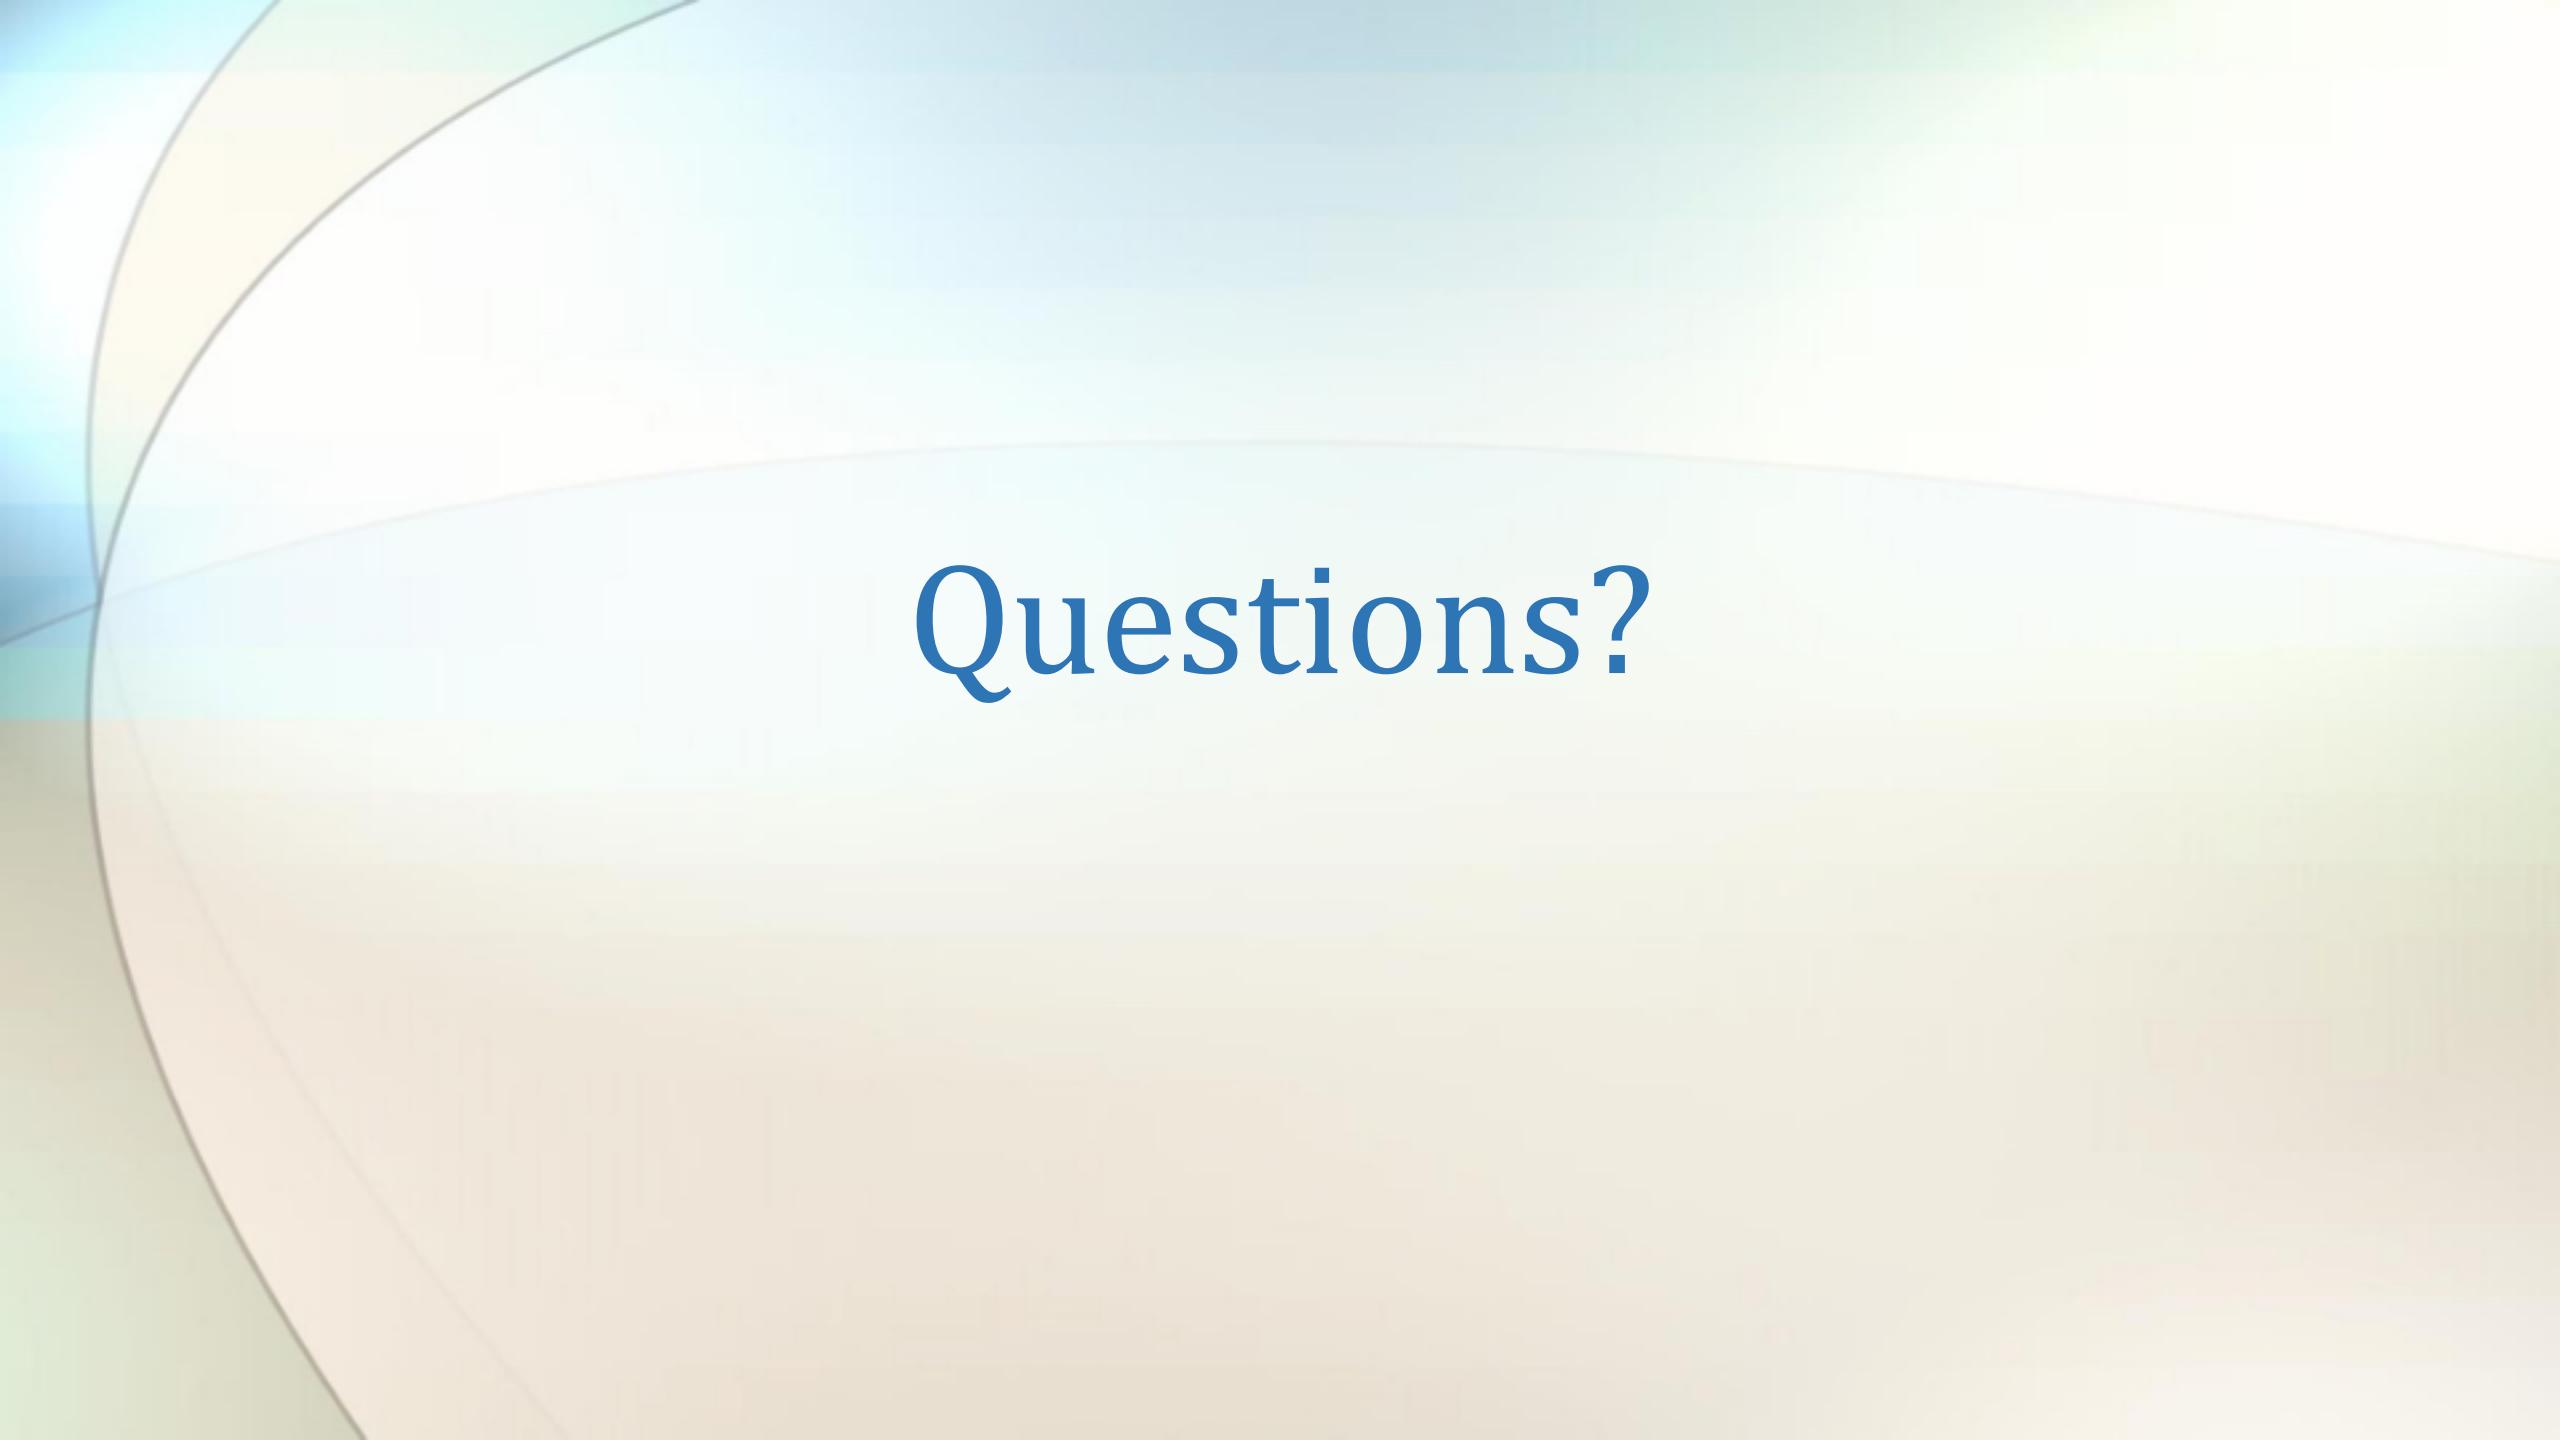
How Well Does It Perform?

```
correctPred = tf.equal(tf.argmax(y,1), tf.argmax(yTruth,1))  
accuracy = tf.reduce_mean(tf.cast(correctPred, tf.float32))  
resultAcc = sess.run(accuracy, feed_dict=  
    {x: mnist.test.images, yTruth: mnist.test.labels})  
print("Trained Acc: %f" % resultAcc)
```

- Approximately 92% accurate
- YOU'VE DONE IT! YOU'VE DONE DEEP LEARNING!!!
 - Kind of, this was a super small, simple, shallow network.
 - 92% is quite bad on this problem
 - Best systems are around 99.7% accurate (Convolutional Neural Networks).

Going Further

- Two main areas to work on machine learning:
 - Architecture, the better the architecture the better the results
 - More layers
 - “Fatter” layers
 - Intertwining layers
 - Tricks like dropout, dropconnect, regularization, pooling, maxout, etc.
 - Network Building Blocks
 - Convolutional Neural Networks
 - Recurrent Neural Networks
 - Generative Adversarial Neural Networks



Questions?

